



Plan-Space Search

• constraint satisfaction problem

→ using constraint-satisfaction techniques → to get flexible solutions
(e.g. plans with partially ordered actions)

BASIC IDEA

- backward search from goal

- each node of search space

↓
partial plan that
contains flaws *

↓
remove the flaws
by making refinements

↓
if successful → we get a PARTIALLY ORDERED SOLUTION

Definitions

✓ Partially ordered Plan

- partially ordered set of nodes
- each node has an action

✓ Partially ordered solution for a planning problem P

- partially ordered plan Π :
- such that every total ordering of Π is a solution for P

✓ partial plan

- partially ordered set of nodes

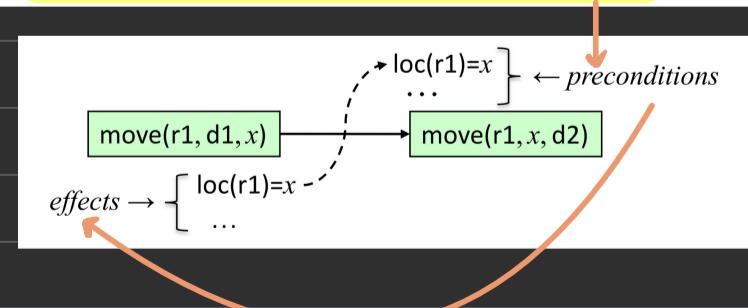
↳ that contain partially instantiated actions

griffonia

- inequality constraints

- causal links (, ^, _, \rightarrow)

- constraint: action a must be the action that establishes action b 's precondition p



* Flaws: 1. Open Goals

Action **b**, precondition **p**

- **p** is an *open goal* if there is no causal link for **p**

↓
Resolve the flaw by creating a causal link

- Find an action **a** (either already in π , or add it to π) that can *establish p*
 - can precede b
 - can have **p** as an effect
- Do substitutions on variables to make **a assert p**
- Add an ordering constraint $a \prec b$
- Create a causal link from **a** to **p** **a makes p true**

(a)
auto to action μπορεί να
1. γινθεί πριν το **b**.
2. να εχει το **p** ως εφέ

eg. 15/36

* Flaws: 2. Threads

(L) causal link from an **effect** of action **a** to a precondition **p** of action **b**

⇒ Action **c** threatens **L** if **c** may come between **a** & **b** & **c** and may affect **p**

- “**c** may come between **a** and **b**” means the plan’s current ordering constraints don’t prevent it
 - plan doesn’t already have $c \prec a$ or $b \prec c$
↳ we have to do this!

to effect λανούμενης καρέ, threat
δραγ μπορεί να μπει οντότητα στο casual
support λανούμενης
action & να κομεί
modify to effect
ΕΚΕΙΝΟ (με replacements
μεταβλητών)

- “**c** may affect **p**” means
 - can substitute values for variables such that **c**’s effects either make **p** true or make **p** false

eg. 16/36

Resolving Threads

- Suppose action **c** threatens a causal link **l** from an effect of action **a** to a precondition **p** of action **b**



3 possible resolves:

1. add a precedence constraint $c < a$
2. add a precedence constraint $b < c$
3. add inequality constraints that prevent c from affecting p

* make sure that c happens before a
(so it can't interfere with a 's effects)

* same with above

* ex. $r_2 \neq r_1$

robot r_2 can't be r_1

- Each of these is applicable iff it doesn't make the plan inconsistent
 - ▶ e.g., 2 isn't applicable if the plan already has $c < b$

PSP Algorithm

$\text{PSP}(\Sigma, \pi)$

loop

```

if  $\text{Flaws}(\pi) = \emptyset$  then return  $\pi$ 
arbitrarily select  $f \in \text{Flaws}(\pi)$ 
 $R \leftarrow \{\text{all feasible resolvers for } f\}$ 
if  $R = \emptyset$  then return failure
nondeterministically choose  $\rho \in R$ 
 $\pi \leftarrow \rho(\pi)$ 
return  $\pi$ 

```

PSP Algorithm:

loop:

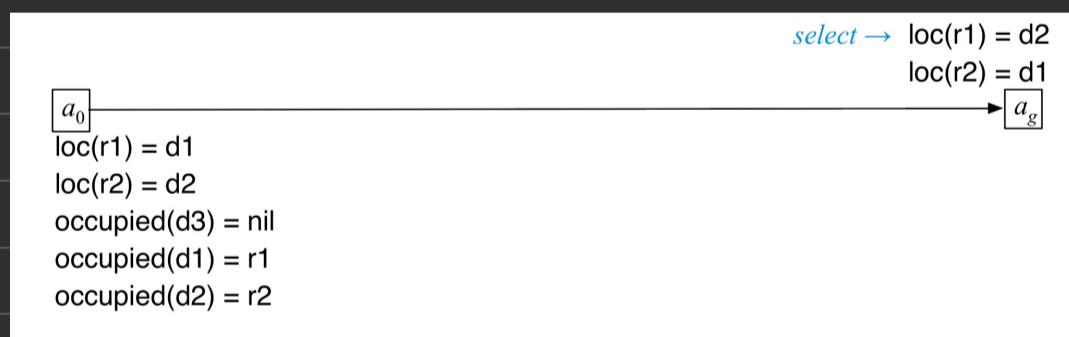
```

av δεν υπάρχουν flaws return the plan
else
  select  $f$  from flaws
   $R \leftarrow$  ολα τα resolvers για το  $f$  που
    είναι possible να συντηνούν
  av δεν υπάρχουν resolvers → Failure
  chose  $\rho \in R$ 
  → return  $\pi \leftarrow \rho(\pi)$ 

```

Full example explanation

- 2 open goals
- no threats



select loc(r1) = d2

open goal

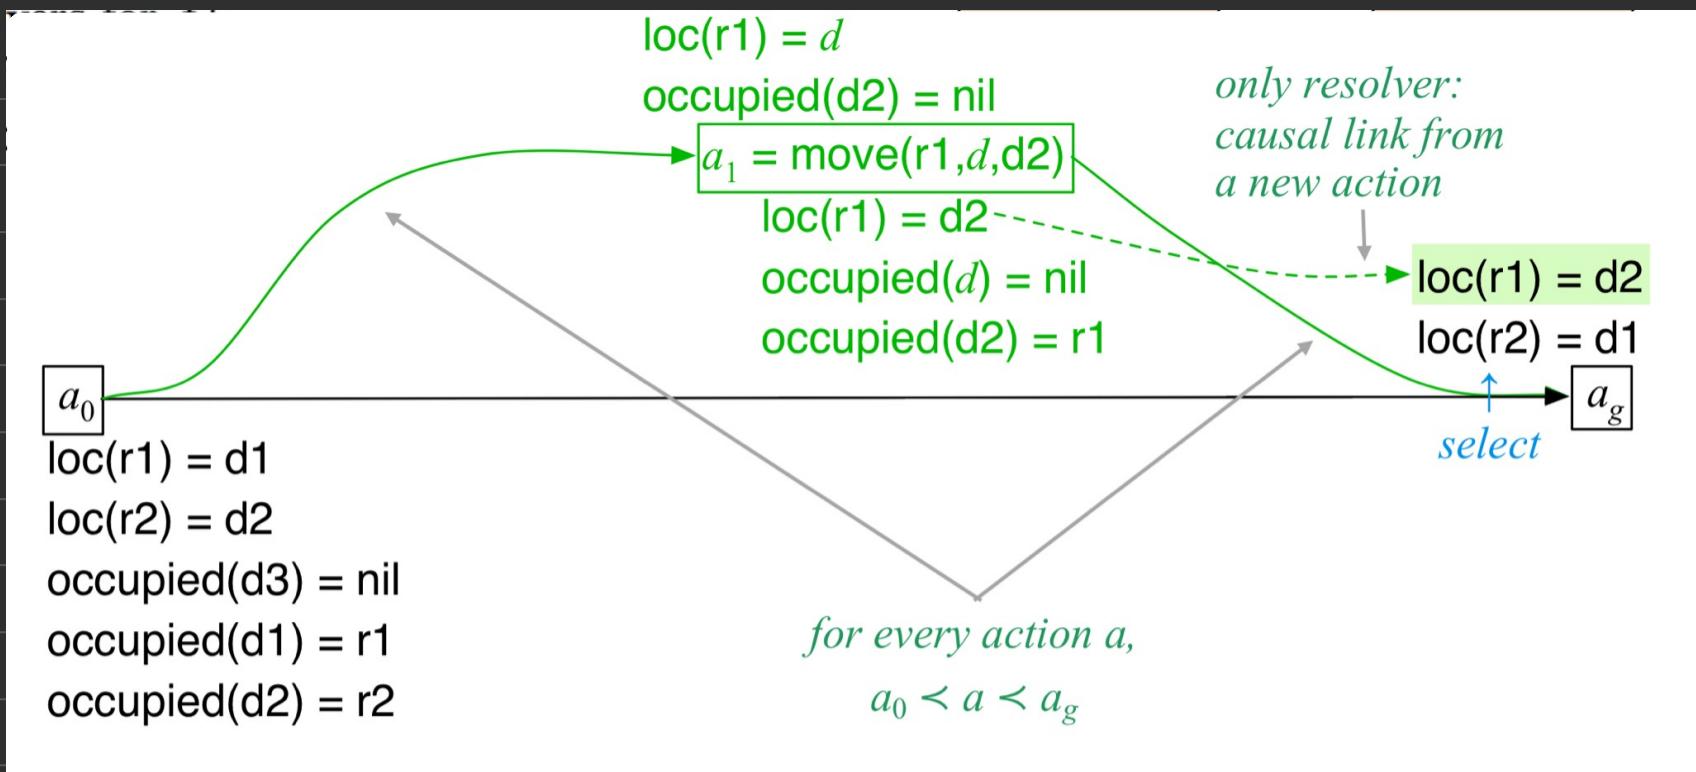
- Find an action a (either already in π , or add it to π) that can establish p
 - can precede b
 - can have p as an effect
- Do substitutions on variables to make a assert p
- Add an ordering constraint $a < b$
- Create a causal link from a to p

action a : πενει να εγονωσει το p
 δηλωση τα effect ons να είναι το p

$a = \text{move}(r1, d, d2) \xrightarrow{\text{effect}} \text{loc}(r1) = d2$

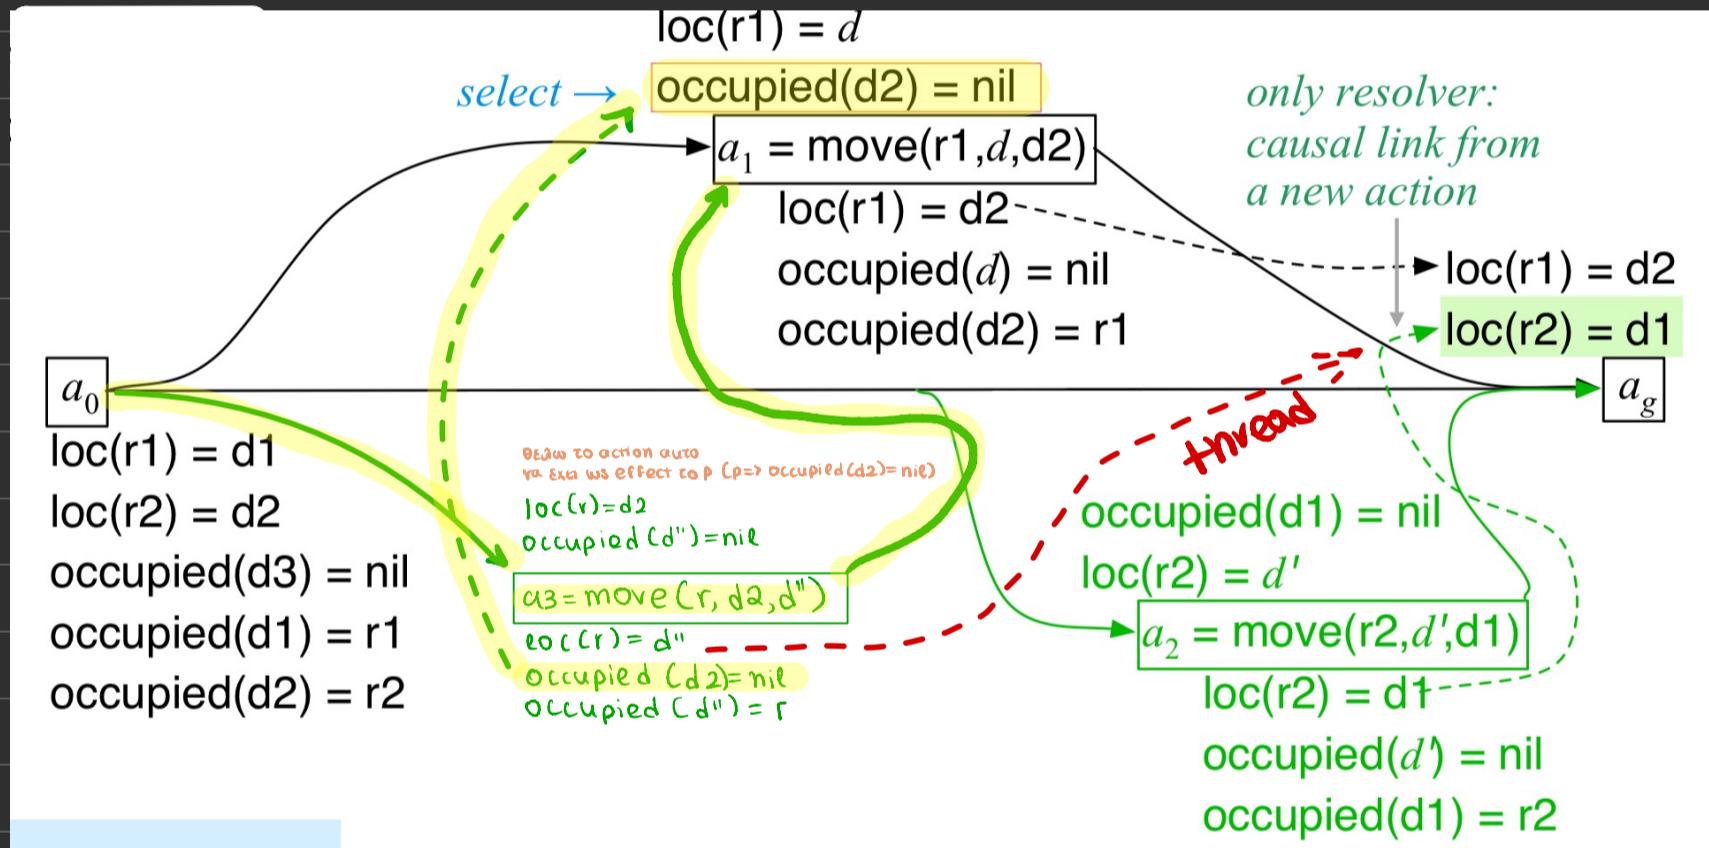
$p \Rightarrow \text{loc}(r1) = d2$

$a \Rightarrow \text{move}(r1, d, d2)$



- 4 open goals
 - no threats
 - $\Rightarrow p \Rightarrow loc(r2) = d1$
- }
- 5 open goals
 - 1 threat

now we select this (bc it is a p with no causal link)



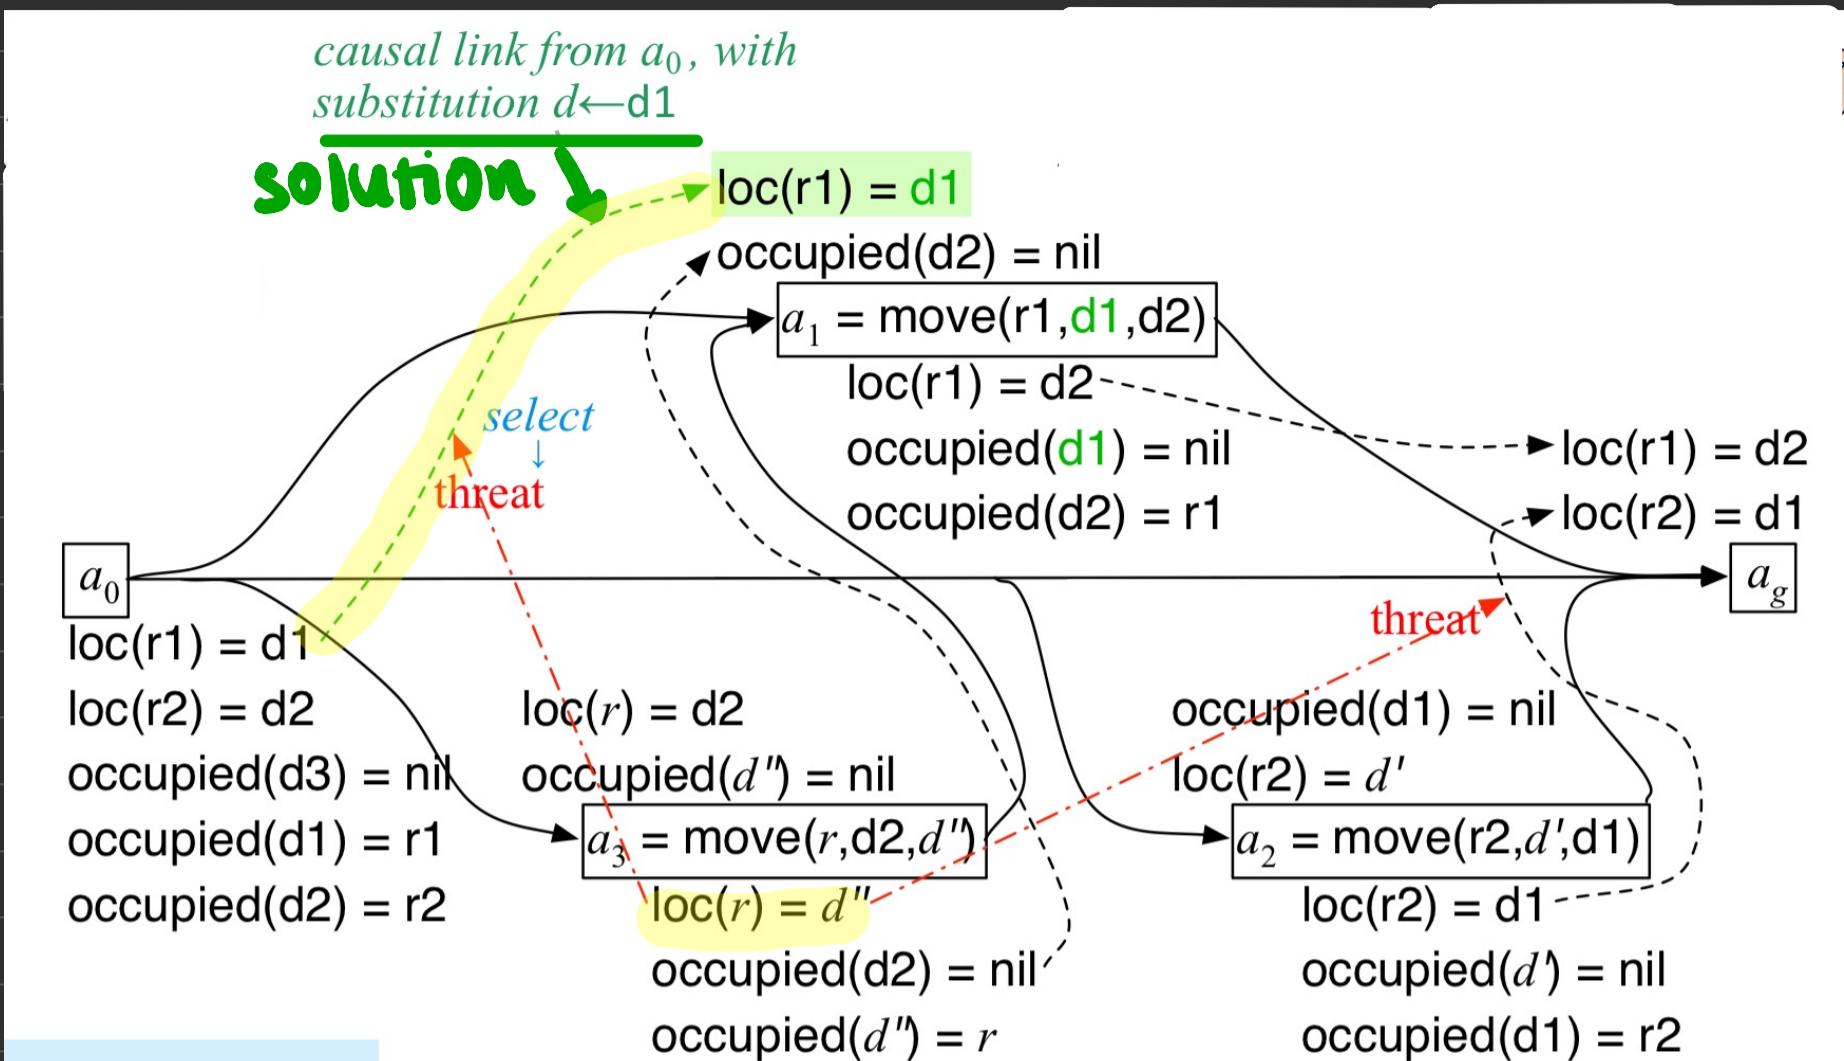
now action a3 node overlaps w/ a2 b/c ag => threat

threat

- ▶ “ c may come between a and b ” means the plan’s current ordering constraints don’t prevent it
 - plan doesn’t already have $c < a$ or $b < c$

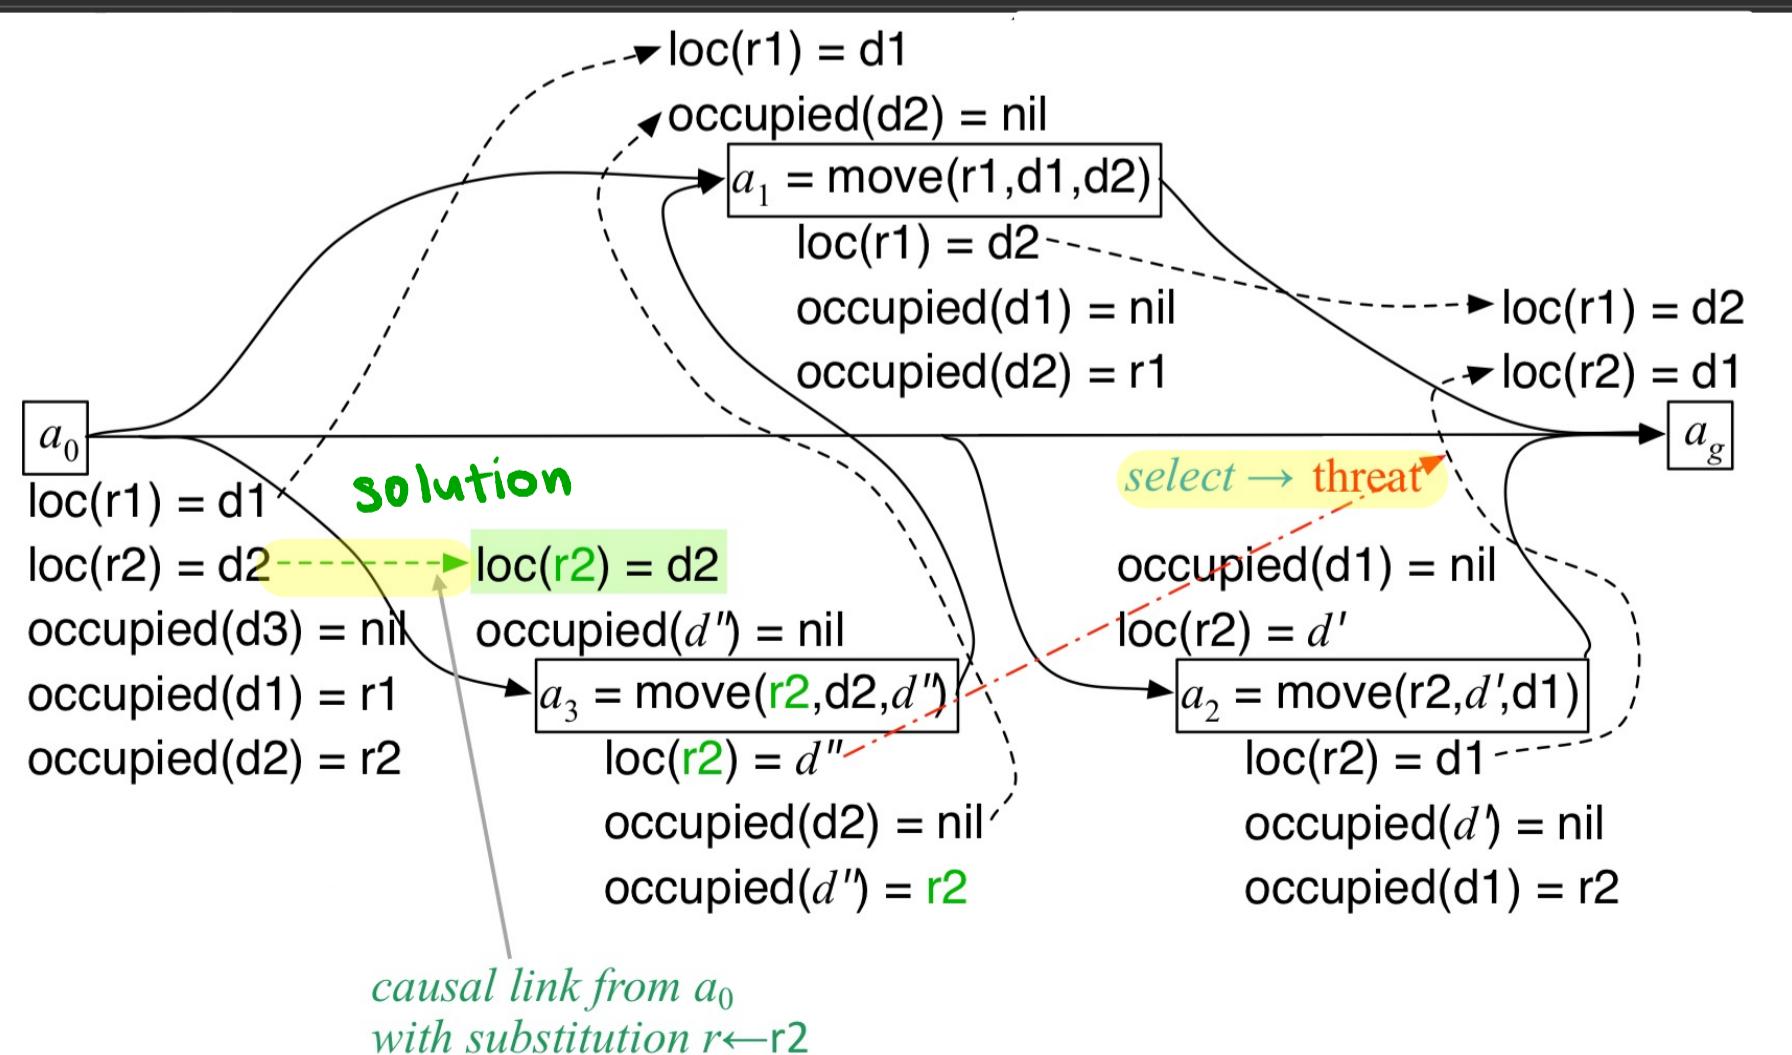
allows c to do that ↴
↳ we have to do this!
- ▶ “ c may affect p ” means
 - can substitute values for variables such that c ’s effects either make p true or make p false ↴

- 4 open goals
- 2 threats

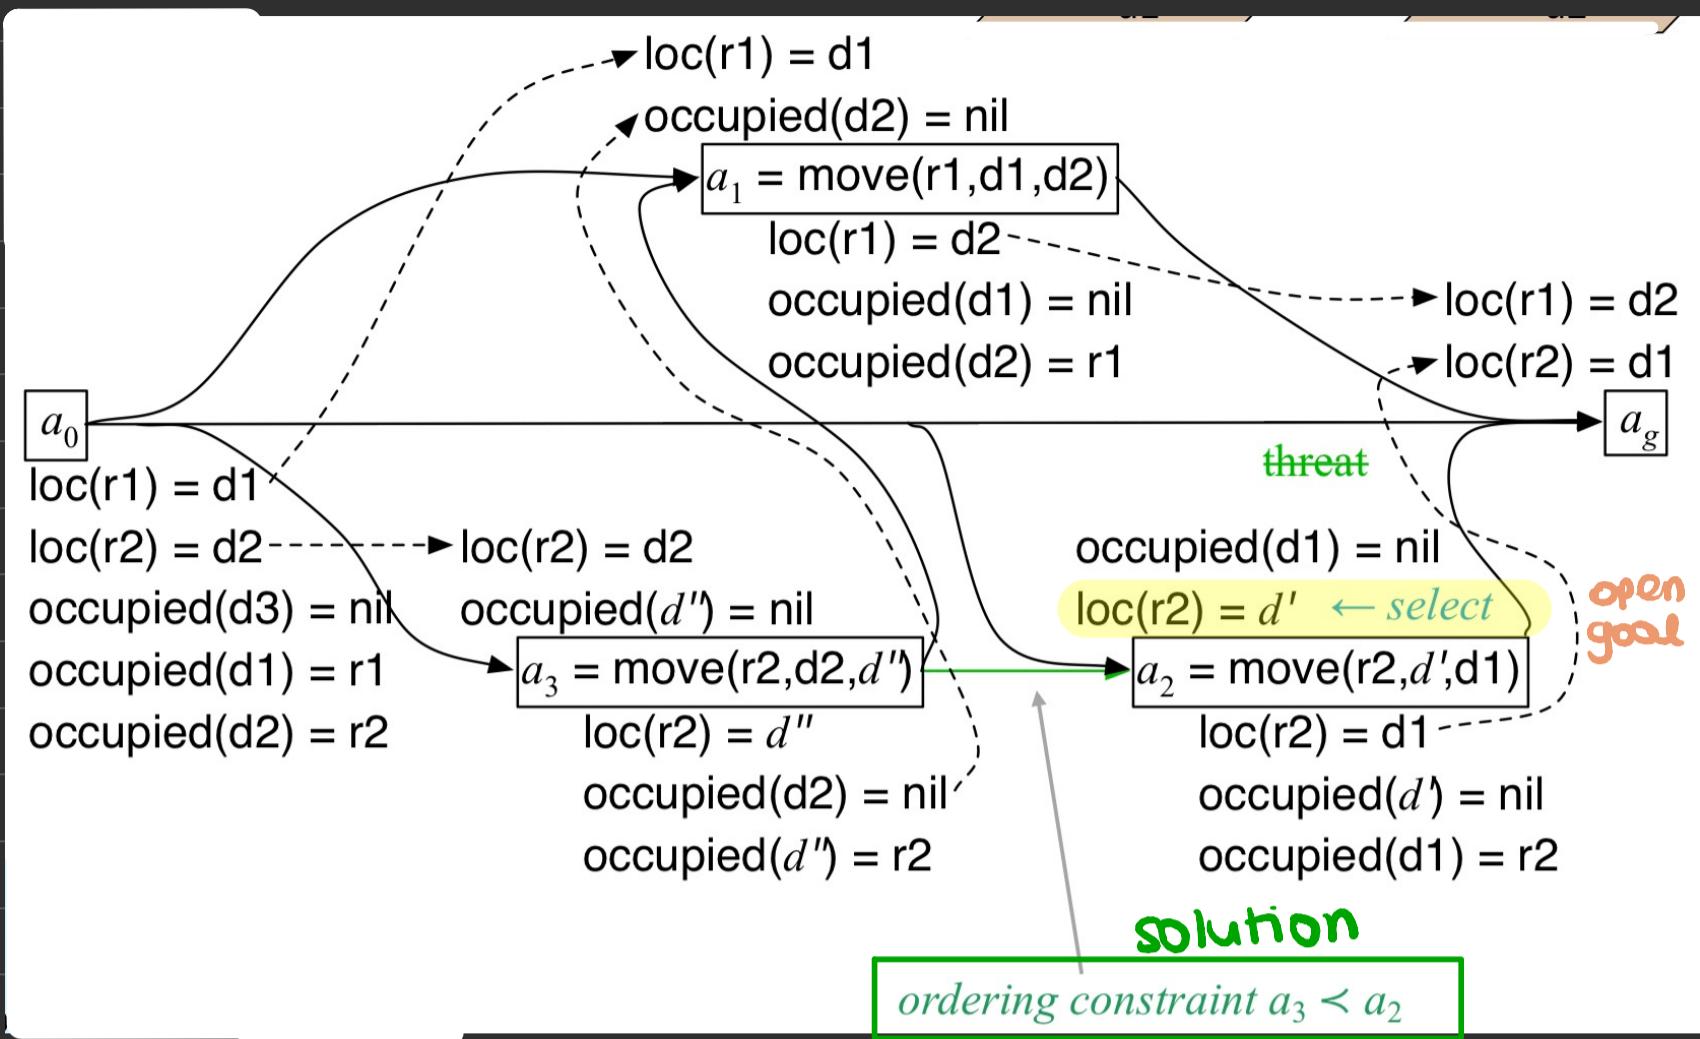


- 4 open goals
- 1 threat

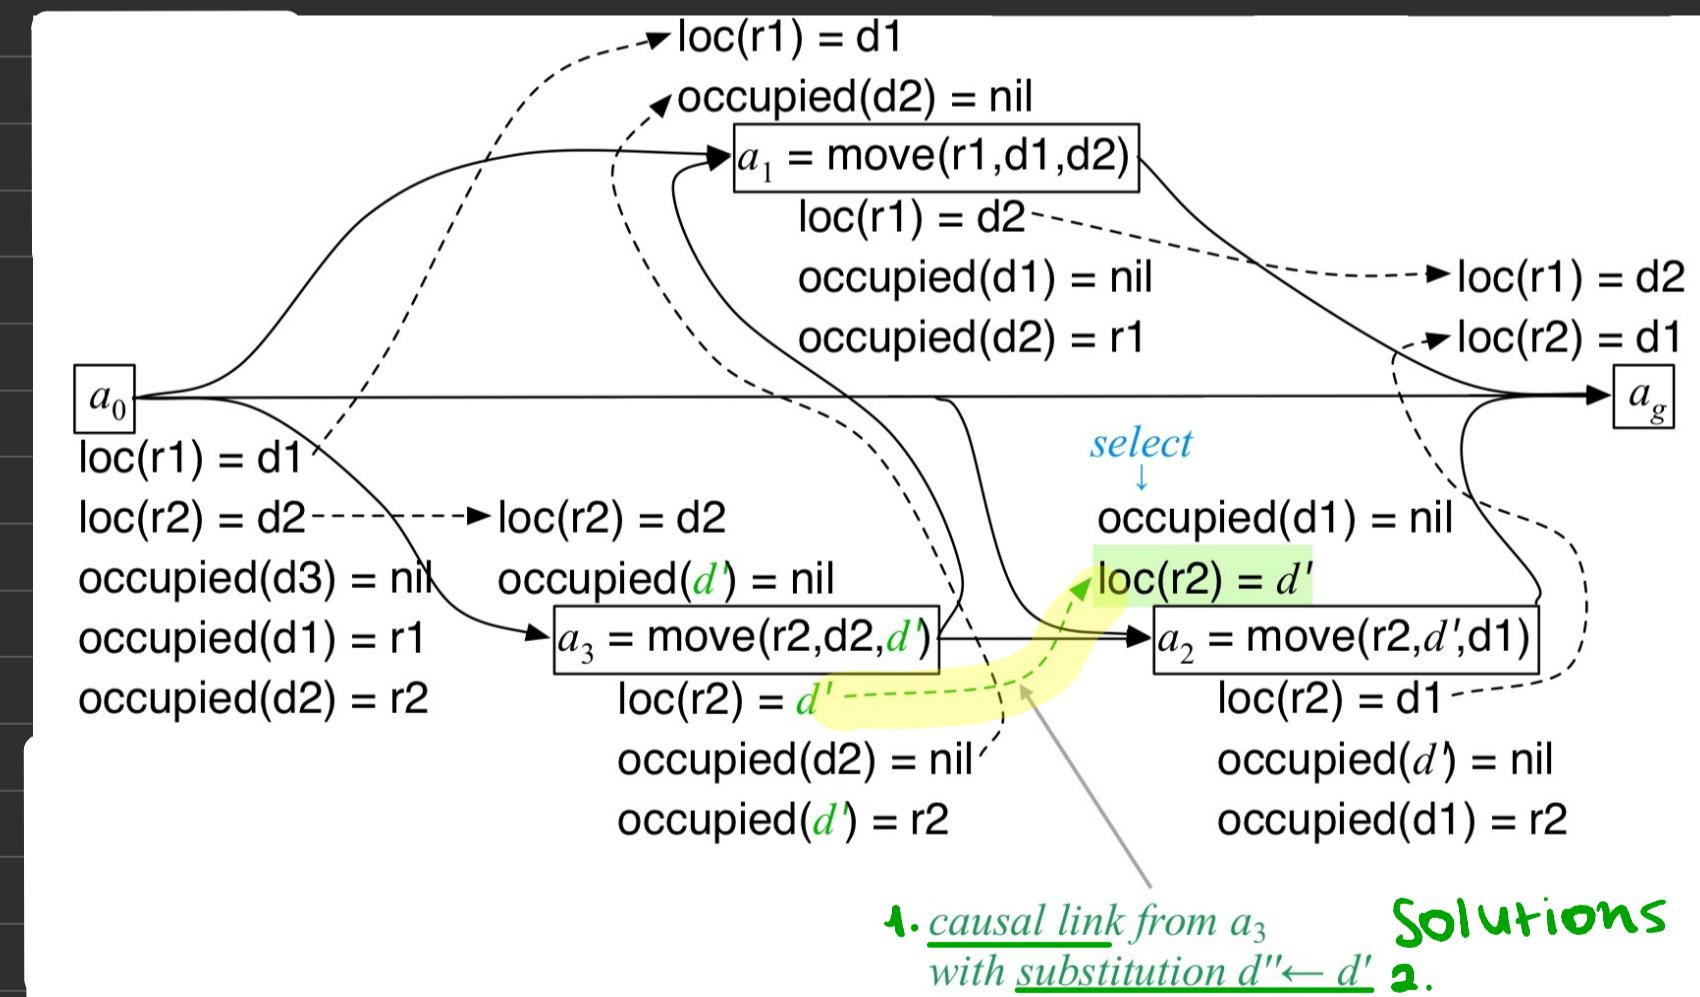
~~constraint : $r \neq r'$~~



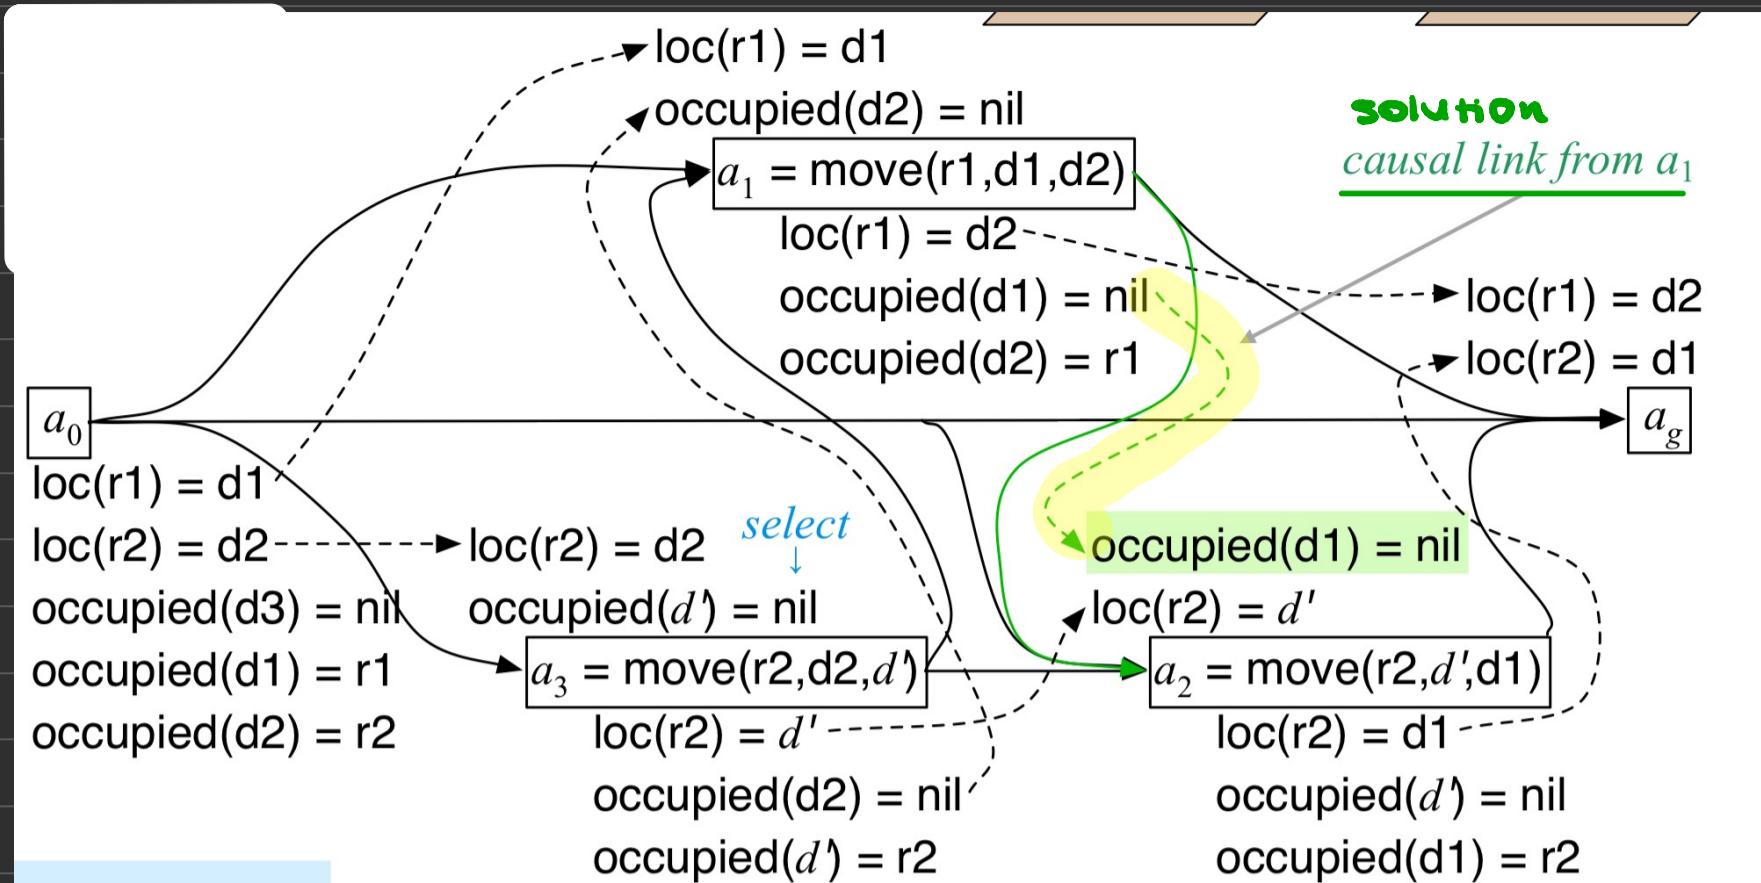
- 3 open goals
- no threats



- 2 open goals
- no threats



- 1 open goal
- no threats



last step!

- no open goals / threats

PSP(Σ, π)

loop

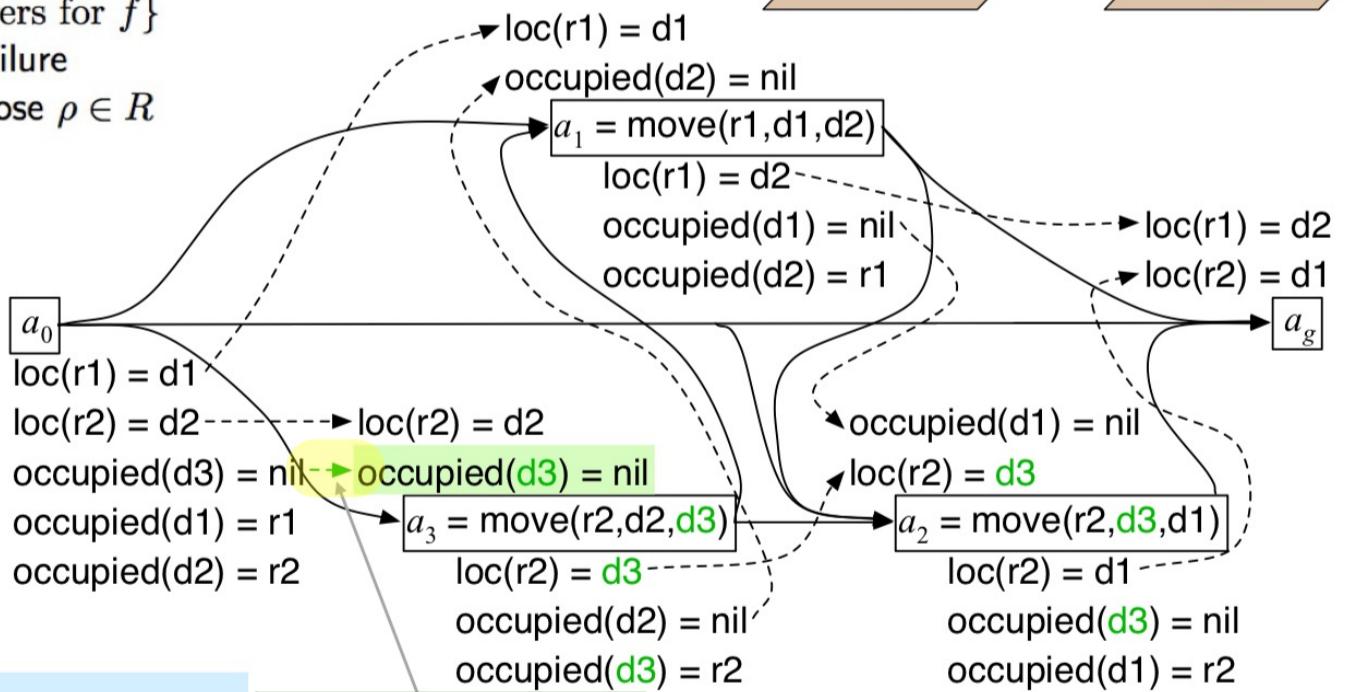
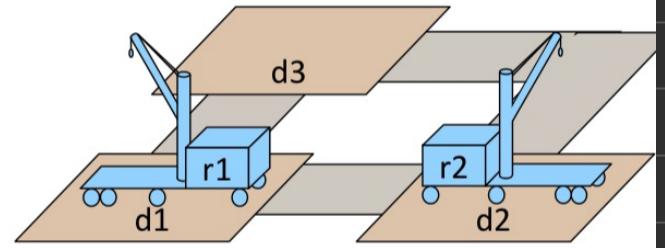
```

    if  $Flaws(\pi) = \emptyset$  then return  $\pi$ 
    arbitrarily select  $f \in Flaws(\pi)$ 
     $R \leftarrow \{\text{all feasible resolvers for } f\}$ 
    if  $R = \emptyset$  then return failure
    nondeterministically choose  $\rho \in R$ 
     $\pi \leftarrow \rho(\pi)$ 
return  $\pi$ 
```

- no open goals
- no threats
- we're done

$move(r; d, d')$
 pre: $loc(r) = d$, $occupied(d') = nil$
 eff: $loc(r) \leftarrow d'$, $occupied(d') = r$, $occupied(d) = nil$

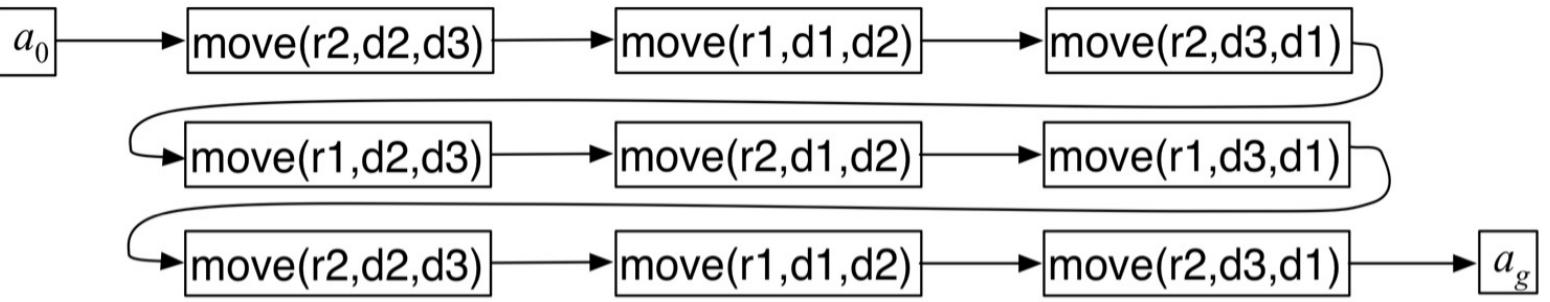
PSP Algorithm



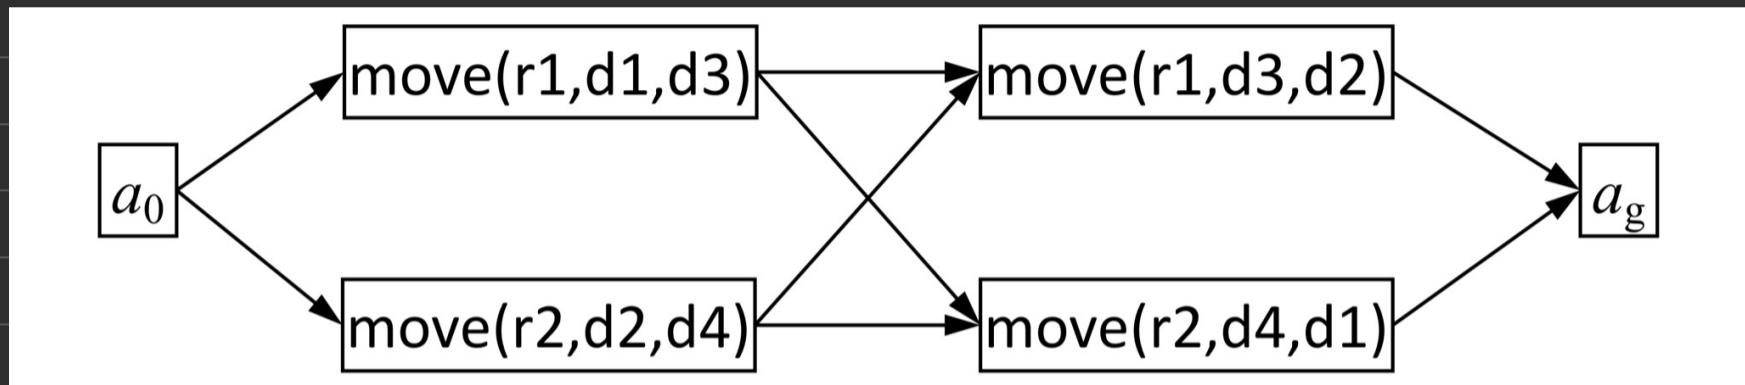
4. causal link from a_0
 9. with substitution $d' \leftarrow d3$

solution

Solutions

- The solution we found: 
- Another: 
- Infinitely many others

- Add another location to the initial state
- This time, PSP also can return PARTIALLY-ORDERED solutions



Selecting a flaw

selecting a flaw to resolve in PSP

≈
Selecting a variable to instantiate in a CSP constraint satisfaction problem

AND-branch in both cases

→ Fewest Alternatives First (FAF)
- select flaw w fewest resolvers

⇒ MRV (min remaining values)
heuristic for CSPs.

choosing a resolver

choosing a resolver for a flaw

≈
assigning a value to a variable in CSP

OR-branch in both cases

-Least Constraining Resolver (LCR)

→ prefer resolver that rules out the fewest resolvers
for the other flows

-LCV (Least Constraining Value) heuristic for CSRs.

problem in PSP

* keeps adding new actions
forever

this might work ANA (Avoid New actions) heuristic

- prefer resolvers that
don't add new actions

- use LCR as tie-breaker

But ::

- Problem: ANA will prefer these two choices:
 - ▶ For $\text{loc}(r1)=d$ in a_1 , use action a_0 with substitution $d \leftarrow d1$
 - ▶ For $\text{loc}(r2)=d'$ in a_2 , use action a_0 with substitution $d' \leftarrow d2$
 - ▶ $a_1 = \text{move}(r1, d1, d2); a_2 = \text{move}(r2, d2, d1)$
 - Makes the problem unsolvable \Rightarrow need to backtrack
- Perhaps use ANA anyway?

Summary

• 2.5 Plan-Space Search

▶ Definitions

- Partially ordered plans and solutions
- partial plans
- causal links

▶ flaws:

- open goals
- threats

▶ resolvers

▶ PSP algorithm

- long example
- brief discussion of node-selection heuristics, pruning techniques

Planning & Acting

Planning

Prediction / Search

- Search over predicted states, possible organizations of tasks and actions
- Uses *descriptive* models (e.g., PDDL)
 - predict what the actions will do*
 - don't include instructions for performing it*

Acting

Performing

- Dynamic, unpredictable, partially observable environment
 - Adapt to context, react to events
- Uses *operational* models
 - instructions telling *how to perform the tasks*
 - usually *hierarchical*

REPRESENTATION

Tasks & Methods

Task → activity for an actor to perform
 taskname (arg₁, ...)

foreach task → 1 or more: refinement methods

↓
 operational models
 telling how to perform the task

methodName(arguments)
 tasks
 preconditions
 body

method-name(arg₁, ..., arg_k)
 task: task-identifier
 pre: test
 body:
a program
 ↓

- assignment statements
- control constructs:
 - if-then-else, while, ...
- tasks
 - can extend this to include events, goals
- commands to the execution platform

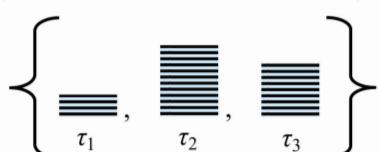
ACTING (RAE)

RAE [Refinement Acting Engine]

↓
 performs multiple tasks in parallel → reactive / no lookahead

↓
 for each task or event τ , a refinement stack → execution stack

- Agenda = {all current refinement stacks}



- Refinement stack for a task τ
 \Leftrightarrow current path in RAE's search tree for τ

refinement tree

procedure RAE:

loop:
 for every new external task or event τ do
 choose a method instance m for τ
 create a refinement stack for τ, m
 add the stack to Agenda
 for each stack σ in Agenda
 call Progress(σ)
 if σ is finished then remove it

Extensions to RAE

- Methods for events
 - e.g., an emergency
 - Methods for goals
 - special kind of task: achievegoal
 - sets up a monitor to see if the goal has been achieved
 - Concurrent subtasks

FULL explained example:
 10 - 23 slides

Pages 9-10: Introduction and Problem Setup

- Environment: The problem involves two robots (r_1, r_2) in a partially observable environment with containers (c_1, c_2) and locations ($loc_0, loc_1, loc_2, loc_3, loc_4$).
- Tasks: The primary task for the robots is to fetch containers.
- Commands: Robots can move, perceive their surroundings, take containers, and put them down.

Page 11: Fetch Task and Methods

- Tasks and Methods: Two methods for the fetch task are described:
- m-fetch1: Used when the container's location is unknown. The robot moves to unviewed locations and perceives them until it finds the container.
- m-fetch2: Used when the container's location is known. The robot directly moves to the location and takes the container.

Pages 12-13: Example Execution

- RAE Loop: The Refinement Acting Engine (RAE) processes tasks by selecting appropriate methods, creating refinement stacks, and progressing through these stacks.
- Example: The task $fetch(r_0, c_2)$ is executed with method m-fetch1 as the container location is initially unknown. The robot perceives various locations to find the container.

Pages 14-15: Refinement Tree and Execution

- Refinement Tree: Shows the hierarchical structure of tasks and methods. The fetch task is decomposed into smaller actions like moving and perceiving.
- Execution: The robot executes the steps in m-fetch1, moving to locations and perceiving them until it finds the container and takes it.

Pages 16-17: Progressing and Handling Failures

- Progress Function: Manages the execution of tasks by checking if the current method step is a command, executing it, and updating the state.
- Handling Failures: If a method step fails (e.g., sensor failure), the RAE retries the task using a different method or a different refinement stack.

Pages 18-19: Example with Sensor Failure

- Sensor Failure Scenario: If the robot's sensor fails while perceiving a location, the task is retried with a different method or by re-executing the current method with different parameters.
- Progress Update: The state is updated accordingly, and the RAE continues processing until the task is successfully completed or all options are exhausted.

Pages 20-21: Simulation and Monte Carlo Rollouts

- Simulating Commands: To choose the best method instance, multiple simulated executions (Monte Carlo rollouts) are performed, and the method with the highest expected utility is selected.
- Expected Utility: Each method's expected utility is estimated based on the simulated outcomes, helping the RAE make informed decisions.

Page 22: Monte Carlo Rollout Details

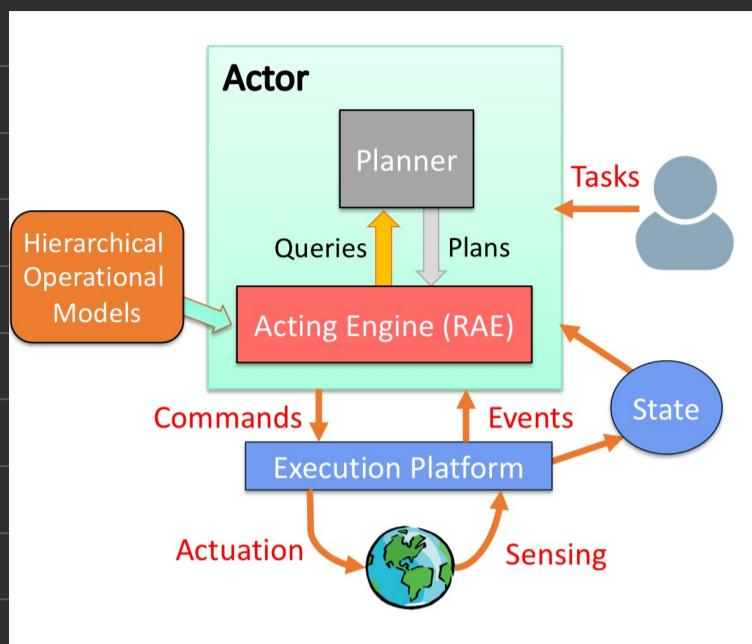
- Rollout Process: Each call to the simulation involves a sequence of actions and their outcomes. The process is repeated multiple times to gather statistical data.
- Multi-arm Bandit Problem: The exploration vs. exploitation dilemma is addressed using algorithms like UCB (Upper Confidence Bound) to balance between trying known successful actions and exploring new ones.

Page 23: Backtracking and Comparison to Search Algorithms

- Backtracking Search: The RAE's retry mechanism resembles a backtracking search, where failed attempts are revisited with alternative methods or parameters.
- Comparison: The approach is compared to classical search algorithms, highlighting the RAE's ability to dynamically adjust to failures and changing environments.

This summary provides a high-level overview of the algorithm example, focusing on the primary tasks, methods, and the process of executing and refining tasks in a dynamic environment.

Planning for RAE



```
procedure RAE:  
loop:  
    for every new external task or event  $\tau$  do  
        choose a method instance  $m$  for  $\tau$   
        create a refinement stack for  $\tau, m$   
        add the stack to  $Agenda$   
    for each stack  $\sigma$  in  $Agenda$   
        call  $Progress(\sigma)$   
        if  $\sigma$  is finished then remove it
```

✓ 4 places where RAE & progress choose a method instance for a task



∴ Bad choice may lead to:

- more costly solution
- failure [need to recover / unrecoverable]

∴ Solution:

- call a PLANNER, choose the method instance it suggests

Planning & Acting Integration

To action models over planners even abstractions

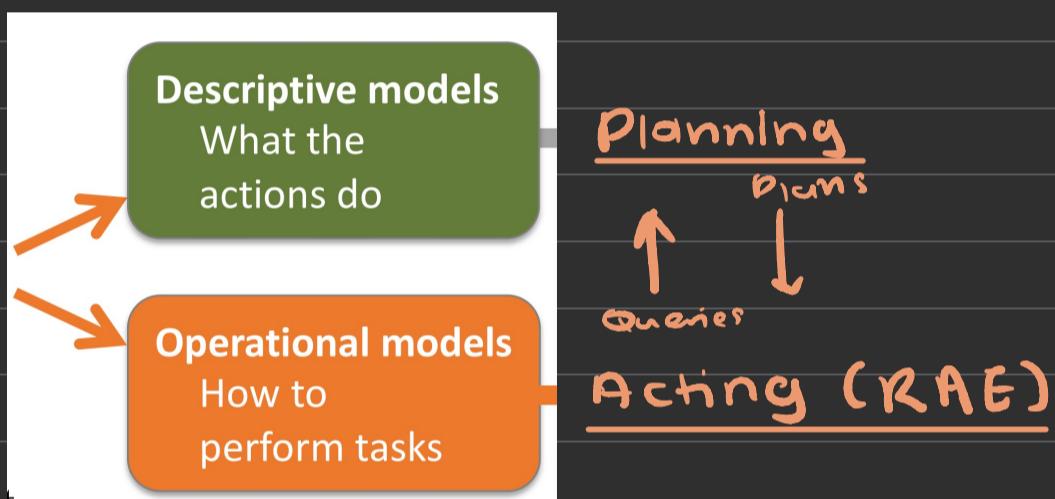
- The planned actions are tasks for the actor to refine

∴ → Consistency problem: how to get action models that describe what the actor will do?

∴ → Solution: ACTOR & PLANNER ^{use} ⇒ SAME REPRESENTATION

- Must be operational; descriptive models too abstract
- Need planning algorithms that can use operational models

consistent?



idea 1

- Planner uses Rae's tasks and refinement methods
- For each of Rae's commands, have a classical action model
- DFS or GBFS search among alternatives to see which works best

SeRPE (Sequential Refinement Planning Engine)

slide 29 + problems slide 30, 31

- Difficult to implement
- Limitations of classical action models

Planning for RAE

```

procedure RAE:
    loop:
        for every new external task or event  $\tau$  do
            choose a method instance  $m$  for  $\tau$ 
            create a refinement stack for  $\tau, m$ 
            add the stack to  $Agenda$ 
        for each stack  $\sigma$  in  $Agenda$ 
            call  $Progress(\sigma)$ 
            if  $\sigma$  is finished then remove it
    
```

idea 2.

- Idea 2: simulation with multithreading or multiprocessing
 - ▶ Run Rae in simulated environment
 - Simulate the commands (see next page)
 - ▶ To choose among method instances, try all of them in parallel
- Planner returns the method instance m having the highest expected utility (\approx least expected cost)

Simulating commands

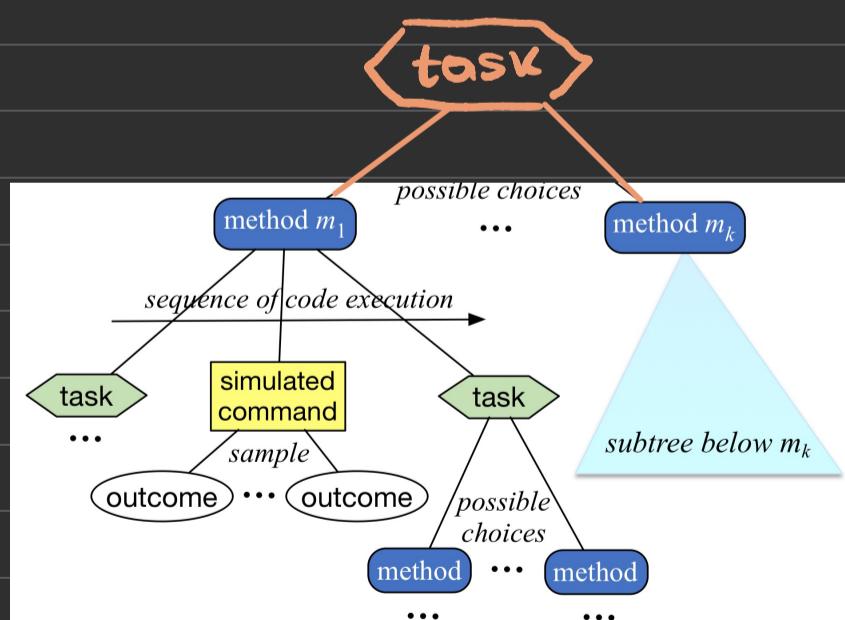
- simplest case:

Probabilistic action template

$a(x_1, \dots, x_k)$
 pre: ...
 (p_1) effects₁: e_{11}, e_{12}, \dots
 ...
 (p_m) effects _{m} : e_{m1}, e_{m2}, \dots

→ choose effects i at random
 with probability p_i and
 use it to update the current state

- More general:
 - ▶ Arbitrary computation, e.g., physics-based simulation
 - ▶ Run the code to get simulated effects



idea 3.

- Idea 3: simulation with Monte Carlo rollouts

- Multiple runs
 - Random choices and outcomes in each run
- Maintain statistics to estimate each choice's expected utility
- Return the method instance m that has the highest estimated utility

Planner

Plan-with-UPOM (task τ):

```
Candidates ← {method instances relevant for  $\tau$ }
for  $i \leftarrow 1$  to  $n$ 
    call UPOM( $\tau$ )
    update estimates of methods' expected utility
return the  $m \in Candidates$  that has
the highest estimated utility
```

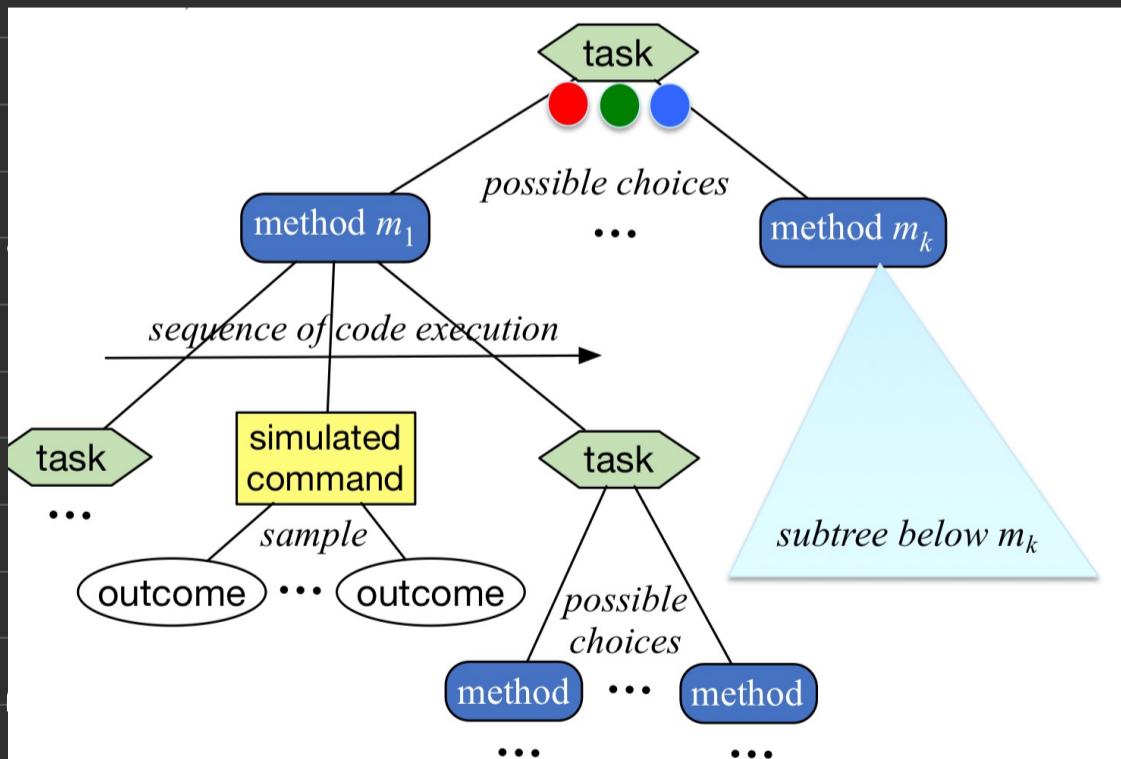
UPOM(τ):

```
choose a method instance  $m$  for  $\tau$ 
create refinement stack  $\sigma$  for  $\tau$  and  $m$ 
loop while Simulate-Progress( $\sigma$ ) ≠ failure
    if  $\sigma$  is completed then return  $(m, utility)$ 
return failure
```

→ each call to UPOM does a Monte Carlo rollout

- simulated execution of RAE on τ

Monte Carlo Rollouts



RAE

Refinement Acting Engine

PI & Act Integr

idea 1: Planner uses RAE's task & refinm. methods



for each RAE's commands



classical
action model

DFS & GBFS chose
which works best → SeRPE

seq.
refin.
plann.
engine

idea 2: simulation via multithreading
&
multiprocessing

Run RAE in simul. environ.

↓
simul. commands

To choose among methods → try them
all in
parallel

Planner returns method having
highest exp. utility

= least expected cost

idea 3: simulation with Monte Carlo Rollouts

↓
multiple runs

↓
random choices & outcomes in each run

↓
statistics → to estimate each
choice's expected
utility

↓
return method instance that has
highest estimated utility

Digression (continued): Monte Carlo Rollouts

- Exploitation vs exploration dilemma:
 - ▶ **Exploitation:** choose an action that has given you high rewards in the past
 - ▶ **Exploration:** choose an action that's less familiar, in hopes that it might produce a higher reward

UCB (Upper Confidence Bound) Algorithm

Assume all rewards are between 0 and 1

For each action a :

- ▶ $r(a)$ = average reward you've gotten from a
- ▶ $n(a)$ = number of times you've tried a
- ▶ $n_t = \sum_a n(a)$
- ▶ $Q(a) = r(a) + \sqrt{2(\ln n_t)/n(a)}$

UCB:

if there are any untried actions:

$\tilde{a} \leftarrow$ any untried action

else:

$\tilde{a} \leftarrow \operatorname{argmax}_a Q(a)$

perform \tilde{a}

update $r(\tilde{a}), n(\tilde{a}), n_t, Q(\tilde{a})$

Theorem (given some assumptions)

As the number of calls to UCB $\rightarrow \infty$
UCB's choice at each call \rightarrow optimal choice

UCT Algorithm

MDP: state space no one-to-one action exec probabilistic outcomes

UCT Algorithm: Monte Carlo Rollouts on a MDP

- At each state s ,
 - ▶ Use UCB to choose an action at random
 - Balances exploration vs exploitation at s
 - ▶ Action's outcome \Rightarrow next state s'

how to use UCT?

call it many times,
return action with
highest expected
utility

Theorem:

as number of calls to UCT $\rightarrow \infty$
choice converges to optimal

by Xander

Convergence

UCT algorithm

- ▶ Monte Carlo rollouts on MDPs
- ▶ Call it many times, choice converges to optimal

UPOM search tree more complicated

- ▶ tasks, method instances, commands, code execution

if no exogenous events

→ can map it to UCT search of a complicated MDP
→ Proof of convergence to optimal

Acting with Planning (RAE + UPOM)

RAE + UPOM

procedure RAE:

loop:

for every new external task or event τ do

* choose a method instance m for τ

create a refinement stack for τ, m

add the stack to *Agenda*

for each stack σ in *Agenda*

call *Progress*(σ)

if σ is finished then remove it

whenever RAE needs to chose a method instance *

=> call Plan-with-UPOM

, use the method instance it returns

Can we use UPOM with Run-Lookahead / Run-Lazy-Lookahead?
slides 43-44-45

comparison

=>

- Rae + UPOM has tighter coupling between planning and acting
 - ▶ works better than Run-Lazy-Lookahead + UPOM'

Learning

Motivation

plan-with-upom gives you called and to RAE, runs online

→ time constraints might not allow complete search

case 1 no time to search at all

- ▶ need a choice function

case 2 enough time to do partial search

- ▶ Receding horizon
 - Cut off search at depth d_{max} or when we run out of time
 - At leaf nodes, use heuristic function to estimate expected utility

Learning algorithms:

Learn Π : learns a choice function

Learn H : learns a heuristic function

Integration with learning

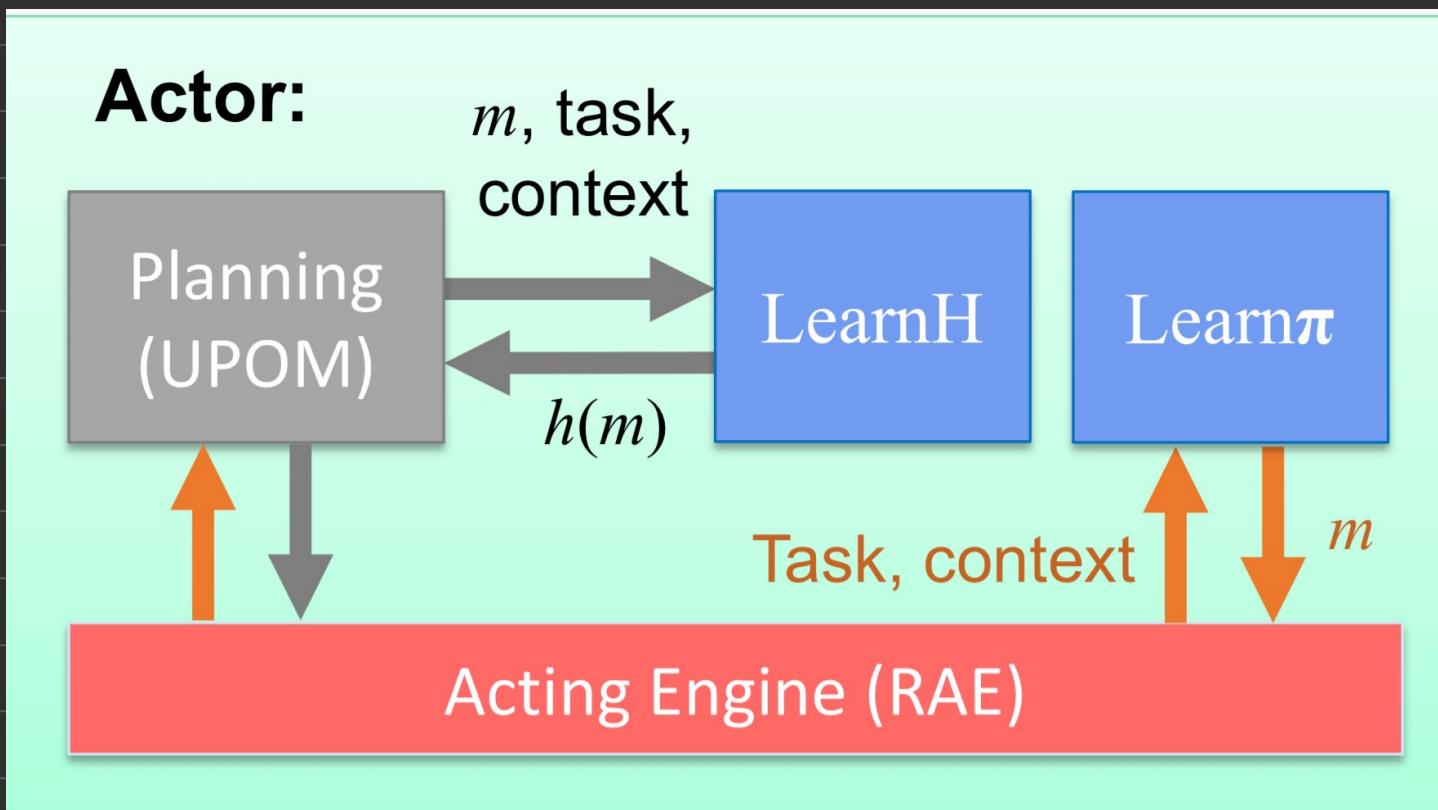
- Gather training data from acting & planning traces of RAE and Plan-with-upom
- Train classifiers

Learn H

learns heuristic function to guide UPOM's search

↓
UPOM can use it to estimate expected utility at leaf nodes

⇒ useful if there isn't enough time to search all the way to the end



Evaluation & Application

Experimental Evaluation

Domain	$ \mathcal{T} $	$ \mathcal{M} $	$ \bar{\mathcal{M}} $	$ \mathcal{A} $	Dynamic events	Dead ends	Sensing	Robot collaboration	Concurrent tasks
S&R	8	16	16	14	✓	✓	✓	✓	✓
Explore	9	17	17	14	✓	✓	✓	✓	✓
Fetch	7	10	10	9	✓	✓	✓	—	✓
Nav	6	9	15	10	✓	—	✓	✓	✓
Deliver	6	6	50	9	✓	✓	—	✓	✓

- Five different domains, different combinations of characteristics
- Evaluation criteria: efficiency (reciprocal of cost), successes vs failures
- Result: Planning and learning help
 - ▶ RAE operates better with UPOM or learning than without
 - ▶ RAE's performance improves with more planning

Prototype Application Software-defined networks

- ▶ Decoupled control and data layers
- ▶ Prone to high-volume, fast-paced online attacks
- ▶ Need automated attack recovery

Prototype Solution using RAE & UPOM

- expert writes recovery procedures as refinements methods

Experimental Results

- Improved efficiency, retry ratio, success ratio, resilience compared to human expert

Summary

- 3.1 Operational models
 - ▶ ξ versus s , tasks, events,
 - ▶ Commands to the execution platform
 - ▶ Extensions to state-variable representation
 - ▶ Refinement method
 - name, task/event, preconditions, body
 - ▶ Example: fetch a container
- 3.2 Refinement Acting Engine (RAE)
 - ▶ Purely reactive: select a method and apply it
 - ▶ Rae: input stream, Candidates, Instances, Agenda, refinement stacks
 - ▶ Progress:
 - command status, nextstep, type of step
 - ▶ Retry: Candidates \ tried
 - comparison to backtracking
 - ▶ Refinement trees
- 3.3 Refinement planning
 - ▶ plan by simulating Rae on a single external task/event/goal
 - ▶ SeRPE uses classical action models
 - ▶ UPOM simulates the actor's commands, does Monte Carlo rollouts
- 3.4 Acting and planning
 - ▶ Rae + UPOM
 - ▶ Comparison: Run-Lazy-Lookahead + UPOM'
 - ▶ A little about learning, experimental evaluation, prototype application
- Open-source Python implementation of Rae and UPOM:
 - ▶ <https://bitbucket.org/sunandita/RAE/>

Refinement based Planning

idea of refining a task into smaller subtasks also used for planning

HTN → hierarchical-task network planning:
process of providing "recipes" for breaking down / refining complex tasks into smaller ones

- ▶ Complex tasks are refined into subtasks using methods.
- ▶ Methods refine tasks in sequence of subtasks.
 - Sorting constraints can be included to generate parallelizable plans.
- ▶ Smaller tasks are modeled as actions from a classic planning domain.

With HTN → planning domains are created with expert knowledge to achieve a solution

- ▶ Whether a solution is reached depends on that knowledge, not just the planner.
- ▶ Planning here is limited to finding the best way to break down the problem.

HTN - Hierarchical Task Network - based planning

For many planning problems → we have ideas of how to find solutions: eg slide 3

- Through HTN we can incorporate this knowledge into a planner.
 - ▶ We will focus on full-order HTN, no order restrictions.
 - ▶ Generates fully ordered, non-parallelizable plans.

Total Order HTN Planning

Ingredients:

- states & actions
- tasks (activities)
- HTN methods (ways to refine tasks)

format of a method:

method-name(args)
Tarea: task-name(args)
Prec: preconditions
Sub: list of subtasks

↓
2 types of Subtasks

/ \

primitives:
name of an action

composed:
needs to be broken down
[= refined]
using methods

- HTN planning domain → pair (Σ, M)

/

Σ: state-variable planning domain (states, actions)

M: methods

- Planning problem:

$P = (\Sigma, M, S_0, T)$

↓

T: a list of tasks
 $\langle t_1, t_2, \dots, t_K \rangle$

solution:

Any executable plan that can be generated by applying:

We apply methods for composed tasks
we apply actions for primitive tasks

- ▶ methods for non-primitive tasks
- ▶ actions for primitive tasks

Planning Algorithm

- depth-first, from left to right
- for each composed tasks, use a method to refine it in subtasks
- for each primitive tasks, apply the corresponding action

e.g. 5-6-7

Integration of hierarchical planning & acting

For planning & acting → necessary that HTN methods can recover from unexpected problems

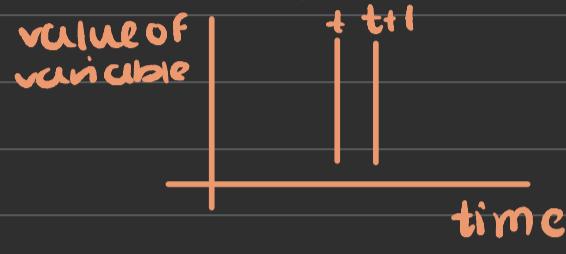
chap4

Notes from video

. RESTRICTIONS OF DURATION IN THE ACTIONS

. WE WORK WITH TIME LINES

↳ they tell us how the value of the variable changes over time



Time line : a pair (T, C)

T temporal assertions C constraints

/ \ change persistence

||
temporal assertions
are the effects of
the actions

notes from professor

* specific parts of the algorithm in the exam:

• working with a timeline to check if it is consistent or secure ✓

• to provide casual support to temporal assertions at a given timeline ✓

• primitive actions & refinement method representation in temporal planning

• ASKS US TO DEFINE A TEMPORAL ACTION OR METHOD

↳ we should know how to define it (setting its temporal assertions, restrictions, etc)

PROPERTIES

consistency

* Timeline is consistent if the 'temporal Assertions' $\rightarrow T$ satisfy all the 'constraints' $\rightarrow C$ that have been set &

no state variable has more than one value at a specific instant of time

- (T, C) is consistent if it has at least one consistent ground instance



(T', C') is consistent if *

Security

* Timeline (T, C) is secure if:

- ✓ it is consistent (at least 1 consistent ground instance)
- ✓ every gr. instance that satisfies the constraints is consistent

* how to make a timeline secure?

- ✓ by adding more constraints



Separation
constraints!

Dojiko!!

Union of Multiple Timelines

- Timelines for k different state variables, all of which are fully ground:
 - $(T_1, C_1), \dots, (T_k, C_k)$
- Union is (T, C) :
 - $T = T_1 \cup \dots \cup T_k$
 - $C = C_1 \cup \dots \cup C_k$
- If every (T_i, C_i) is secure, then $(T_1, C_1) \cup \dots \cup (T_k, C_k)$ is also secure

book omits
this part

Causal Support

what ensures the correct starting value?

action on the temporal assertion



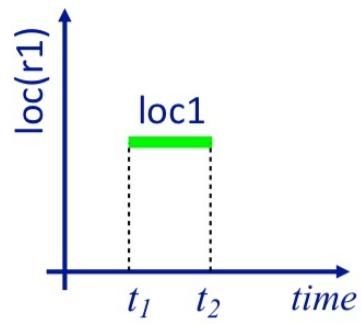
gives causal support to another
temporal assertion

or

it is part of the initial state
(priori causal support)

Causal support

- Consider the assertion $[t_1, t_2] \text{loc}(r1) = \text{loc1}$
 - How did r1 get to loc1 in the first place?
- Temporal Assertions**
- Let α be either $[t_1, t_2] x = v_1$ or $[t_1, t_2] x : (v_1, v_2)$
- Causal support** for α
 - Information saying α is supported *a priori*
 - Or another assertion that produces $x = v_1$ at time t_1
 - $[t_0, t_1] x = v_1$
 - $[t_0, t_1] x : (v_0, v_1)$



- A timeline \mathcal{T} is *causally supported* if every assertion α in \mathcal{T} has a causal support

- Three ways to modify a timeline to add causal support ...

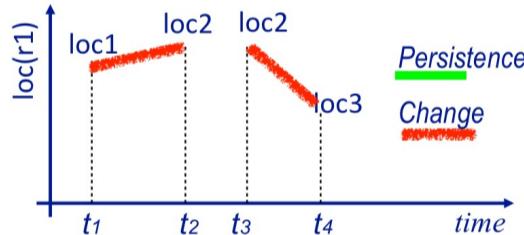
3 ways to modify timeline to add causal support

① Add a persistence assertion

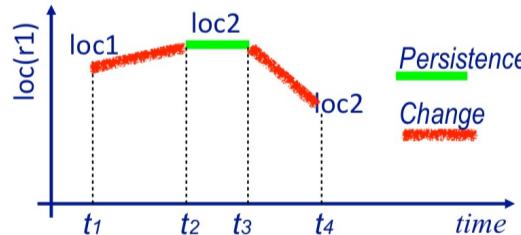
(1) Add a persistence assertion

$$\mathcal{T} = \{[t_1, t_2] \text{loc}(r1):(loc1, loc2), [t_3, t_4] \text{loc}(r1):(loc2, loc3)\}$$

$$C = \{t_1 < t_2 < t_3 < t_4\}$$



- Add $[t_2, t_3] \text{loc}(r1) = \text{loc2}$
 - Supported by the first temporal assertion
 - Supports the second one

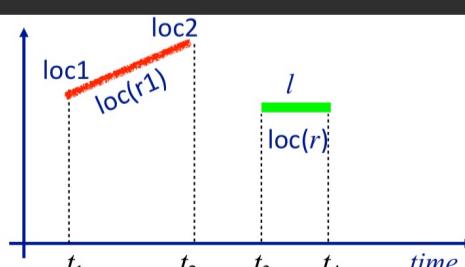


② Add constraints

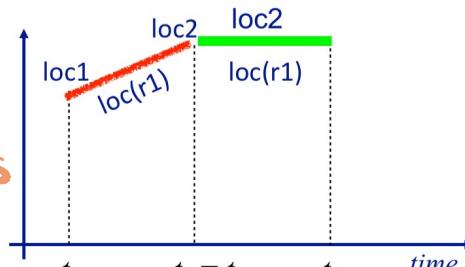
(2) Add constraints

$$\mathcal{T} = \{[t_1, t_2] \text{loc}(r1):(loc1, loc2), [t_3, t_4] \text{loc}(r) = l\}$$

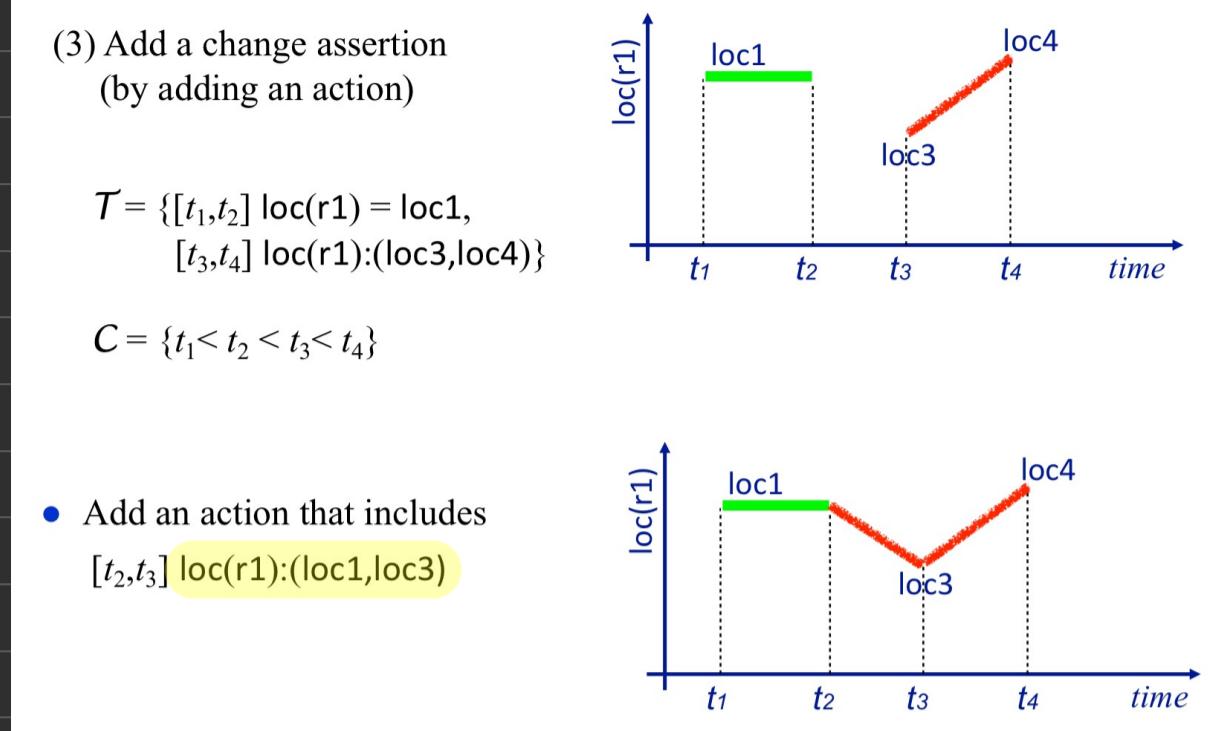
$$C = \{t_1 < t_2, t_3 < t_4\}$$



- Add $t_2 = t_3, r = r1, l = \text{loc2}$
- + assigning specific values to variables



③ Add a change assertion (by adding an action)



!!

the causal support is the essence of the algorithm (later discussed)

How the actions are defined?
How do they include temporal assertions?

Temporal Planning → Planning & Refinement

Actions

/
Primitive

\ can be resolved
through refinement methods
(subdividing them into smaller
actions)

↓
Primitive Tasks (actions)

(head, T, C)

↓
name,
parameters

↳ (T, C) union
of set of timelines

↓ parameters → ts, te

- Always two additional parameters
 - ▶ starting time t_s , ending time t_e
- In each temporal assertion in T ,
 - left endpoint is like a precondition
↔ need for causal support
 - right endpoint is like an effect

T : prec effect

temp.
assert

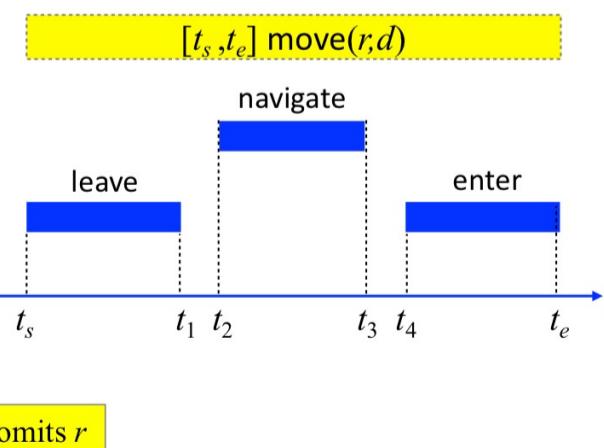
Refinement Methods

Methods & Tasks

how is the tasks refined into primitive actions or other subtasks

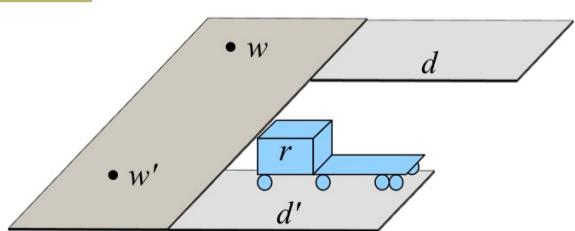
Tasks and Methods

- Task: move robot r to dock d
 - ▶ $[t_s, t_e]$ move(r, d)
adds additional parameters (objects) that are going to appear somewhere in the refinement



- Method:
m-move1(r, d, d', w, w')
 - task: move(r, d)
 - refinement:
- book omits r

- $[t_s, t_1]$ leave(r, d', w')
- $[t_2, t_3]$ navigate(r, w', w)
- $[t_4, t_e]$ enter(r, d, w)



+++

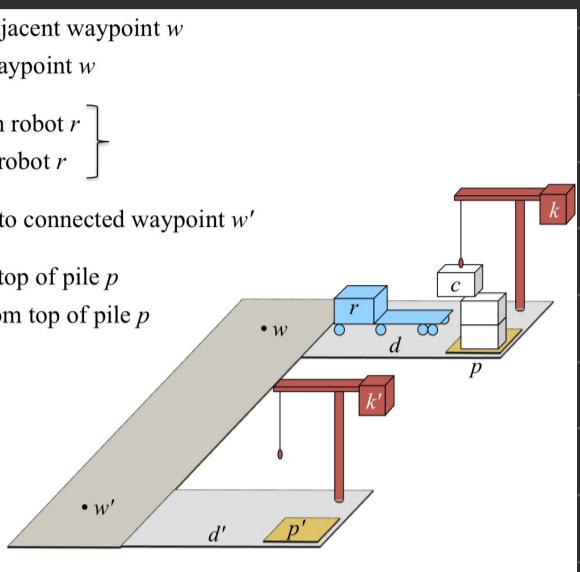
apa 5inv oucia Example

primitive actions

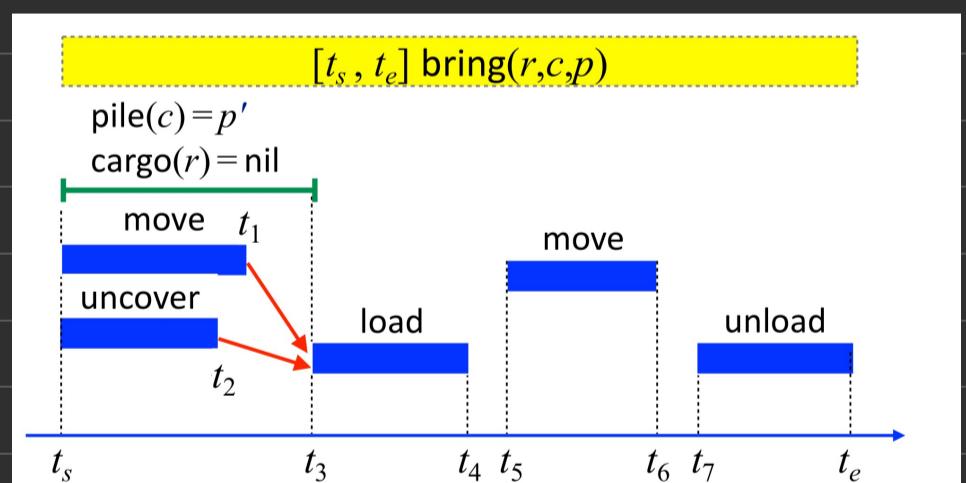
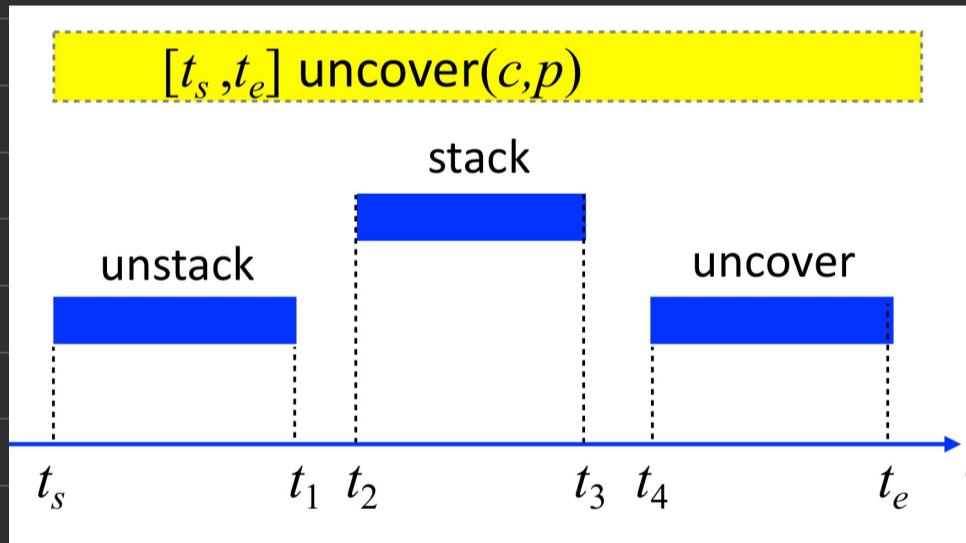
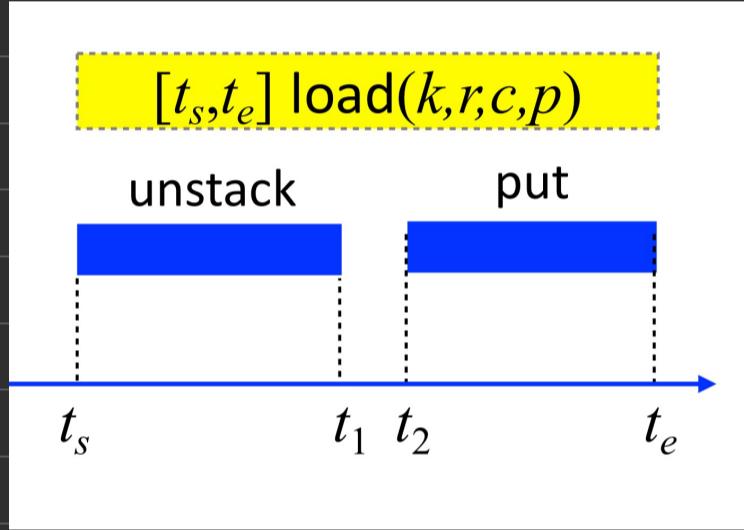
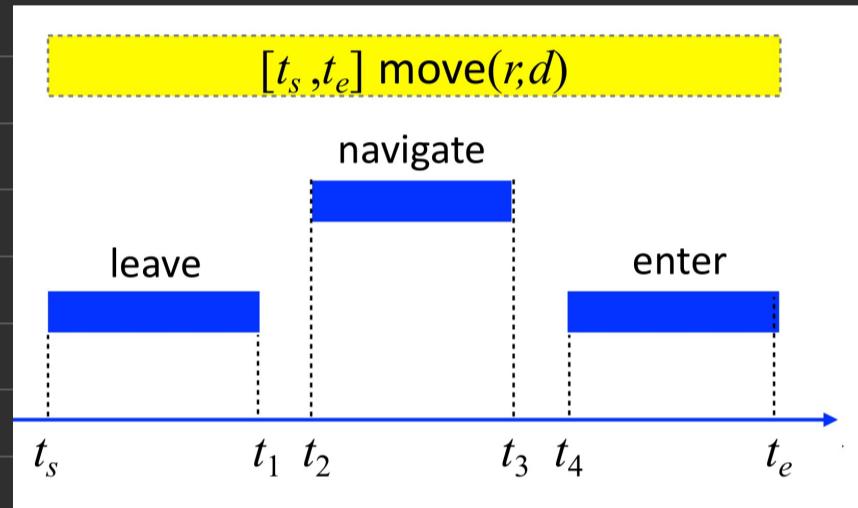
aura:

- leave(r, d, w) robot r leaves dock d to an adjacent waypoint w
- enter(r, d, w) r enters d from an adjacent waypoint w
- take(k, c, r) crane k takes container c from robot r
- put(k, c, r) crane k puts container c onto robot r
- navigate(r, w, w') r navigates from waypoint w to connected waypoint w'
- stack(k, c, p) crane k stacks container c on top of pile p
- unstack(k, c, p) crane k takes a container c from top of pile p

c, c' - containers
 d, d' - loading docks
 k, k' - cranes
 p, p' - piles
 r - robot
 w, w' - waypoints



ταυ μεθόδος του τα χρησιμοποιούν



Summed up

primitive actions $(\text{head}, \mathcal{T}, \mathcal{C})$

head \rightarrow enter (r, d, w)

T \rightarrow assertions:

$[t_s, t_e] \text{ loc}(r) : (w, d)$

$[t_s, t_e] \text{ occupand} : (\text{empty}, r)$

C \rightarrow constraints:

$t_e \leq t_s + \delta_2$

$\text{adjacent}(d, w)$

methods

what includes

m-move (r, d, d', w, w')

task move(r, d)

refinement [list of primitive
actions may have time]
 $[t_s, t_e] \text{ leave}(r, d', w')$...

assertions

$[t_s, t_{s+1}] \text{ loc}(r) = d' // \text{can have current state?}$

constraints

$\text{adjacent}(d, w)$

$t_1 \leq t_2, t_3 \leq t_4$

chronicles

$\varphi = (A, S, T, C)$

- A : temporally qualified tasks
- S : *a priori* supported assertions
- T : temporally qualified assertions
- C : constraints

"a problem to be solved"

constraint management

slide 53

Consistency of C

- C contains two kinds of constraints
 - ▶ Object constraints
 - $\text{loc}(r) \neq l_2, l \in \{\text{loc3}, \text{loc4}\}, r = \text{r1}, o \neq o'$
 - ▶ Temporal constraints
 - $t_1 < t_3, a < t, t < t', a \leq t' - t \leq b$
- Assume object constraints are independent of temporal constraints and vice versa
 - ▶ exclude things like $t < \text{distance}(r, r')$
- Then two separate subproblems
 - ▶ (1) check consistency of object constraints
 - ▶ (2) check consistency of temporal constraints
 - ▶ C is consistent iff both are consistent

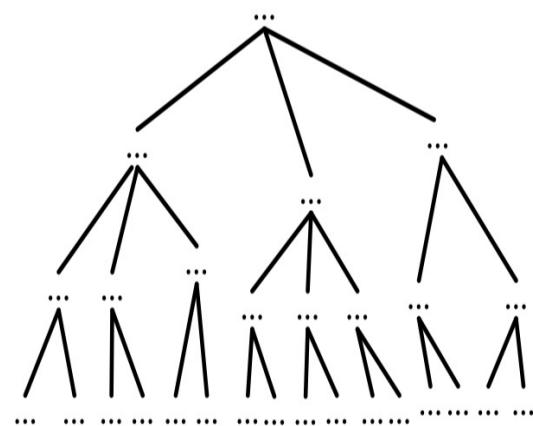
(1)

Object Constraints

- Constraint-satisfaction problem (CSP) – NP-hard
- Can write an algorithm that's *complete* but runs in exponential time
 - If there's an inconsistency, always finds it
 - Might do a lot of pruning, but spend lots of time at each node

- Instead, use a technique that's incomplete but takes *polynomial* time
 - arc consistency, path consistency*
- Detects some inconsistencies but not others
 - Runs much faster, but prunes fewer nodes

*See Russell & Norvig, Artificial Intelligence: A Modern Approach

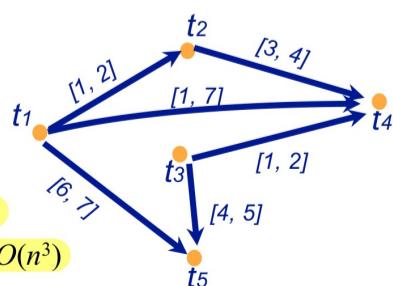


(2)

Time Constraints

To represent time constraints:

- Simple Temporal Networks (STNs)
 - ▶ Networks of constraints on time points
- Synthesize incrementally them starting from ϕ_0
 - ▶ Templan can check time constraints in time $O(n^3)$
- Incrementally instantiated at acting time
- Kept consistent throughout planning and acting



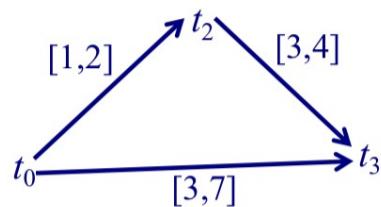
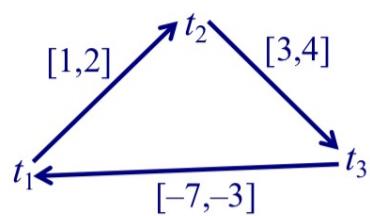
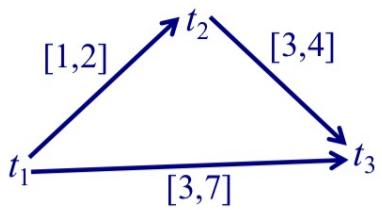
STN: $(V, E) \rightarrow E \subseteq V^2$ set of arcs
set of temporal variables t_1, \dots, t_n

> each arch $[t_i, t_j] \rightarrow$ interval $[a, b]$
represents a constraint $t_j - t_i \in [a, b]$
 $t_i - t_j \in [-b, -a]$

> unary constraints: $t_0 \equiv 0$
 $t_{oi} = [a, b] \Leftrightarrow t_i - 0 \in [a, b]$

Time Constraints

- Simple Temporal Network (STN):
- a pair (V, E) , where
 - $V = \{t_1, \dots, t_n\}$
 - $E \subseteq V^2$ is a set of arcs
- Each arc (t_i, t_j) is labeled with an interval $[a, b]$
 - Represents constraint $t_j - t_i \in [a, b]$
 - Or equivalently, $t_i - t_j \in [-b, -a]$
- Notation: instead of $t_j - t_i \in [a, b]$, write $r_{ij} = [a, b]$
- To represent unary constraints:
 - ▶ Dummy variable $t_0 \equiv 0$
 - ▶ Arc $r_{0i} = [a, b]$ represents $t_i - 0 \in [a, b]$

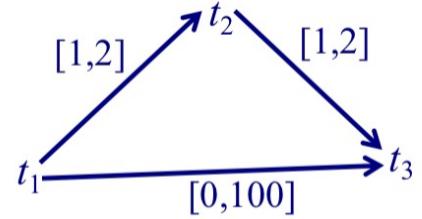


Solution to an STN: va bawm integer values ona time points etci wte oza ra constraints va kavoro iouvzou

↓
consistent STN : has a solution ✓

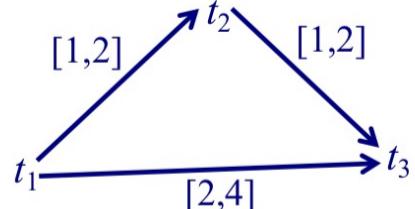
Time Constraints

- Solution to an STN:
 - ▶ any assignment of integer values to the time points such that all the constraints are satisfied
- Consistent STN: has a solution
- Minimal STN:
 - for every arc (t_i, t_j) with label $[a, b]$,
 - for every $t \in [a, b]$,
 - there's at least one solution such that $t_j - t_i = t$
- If we make any of the time intervals shorter, we'll exclude some solutions



- Solutions:

- ▶ $(t_2 - t_1, t_3 - t_2, t_3 - t_1) \in \{(1,1,2), (1,2,3), (2,1,3), (2,2,4)\}$
- ↳ $1 \in t_2 - t_1$
- ↳ $1 \in t_3 - t_2$
- ↳ $2 \in t_3 - t_1$



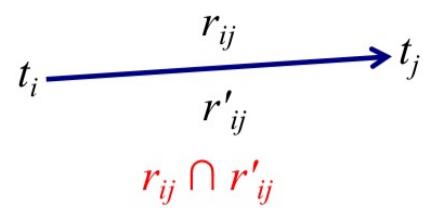
Operations on STNs

- Intersection, \cap

$$t_j - t_i \in r_{ij} = [a_{ij}, b_{ij}]$$

$$t_j - t_i \in r'_{ij} = [a'_{ij}, b'_{ij}]$$

Infer $t_j - t_i \in r_{ij} \cap r'_{ij} = [\max(a_{ij}, a'_{ij}), \min(b_{ij}, b'_{ij})]$



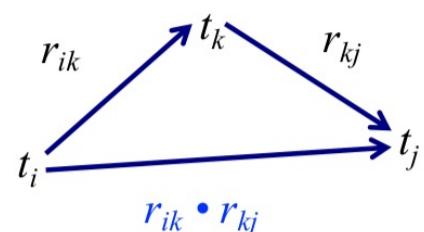
- Composition, \bullet

$$t_k - t_i \in r_{ik} = [a_{ik}, b_{ik}]$$

$$t_j - t_k \in r_{kj} = [a_{kj}, b_{kj}]$$

Infer $t_j - t_i \in r_{ik} \bullet r_{kj} = [a_{ik} + a_{kj}, b_{ik} + b_{kj}]$

Reason: shortest and longest times for the two intervals

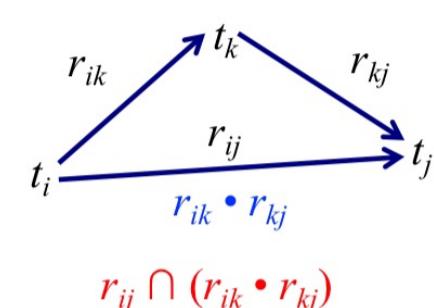


- Consistency checking

► r_{ik}, r_{kj}, r_{ij} are consistent only if $r_{ij} \cap (r_{ik} \bullet r_{kj}) \neq \emptyset$ not empty

- Special case for networks with just three nodes

► Consistent iff $r_{ij} \cap (r_{ik} \bullet r_{kj}) \neq \emptyset$



$$r_{ij} \cap (r_{ik} \bullet r_{kj}) \neq \emptyset$$

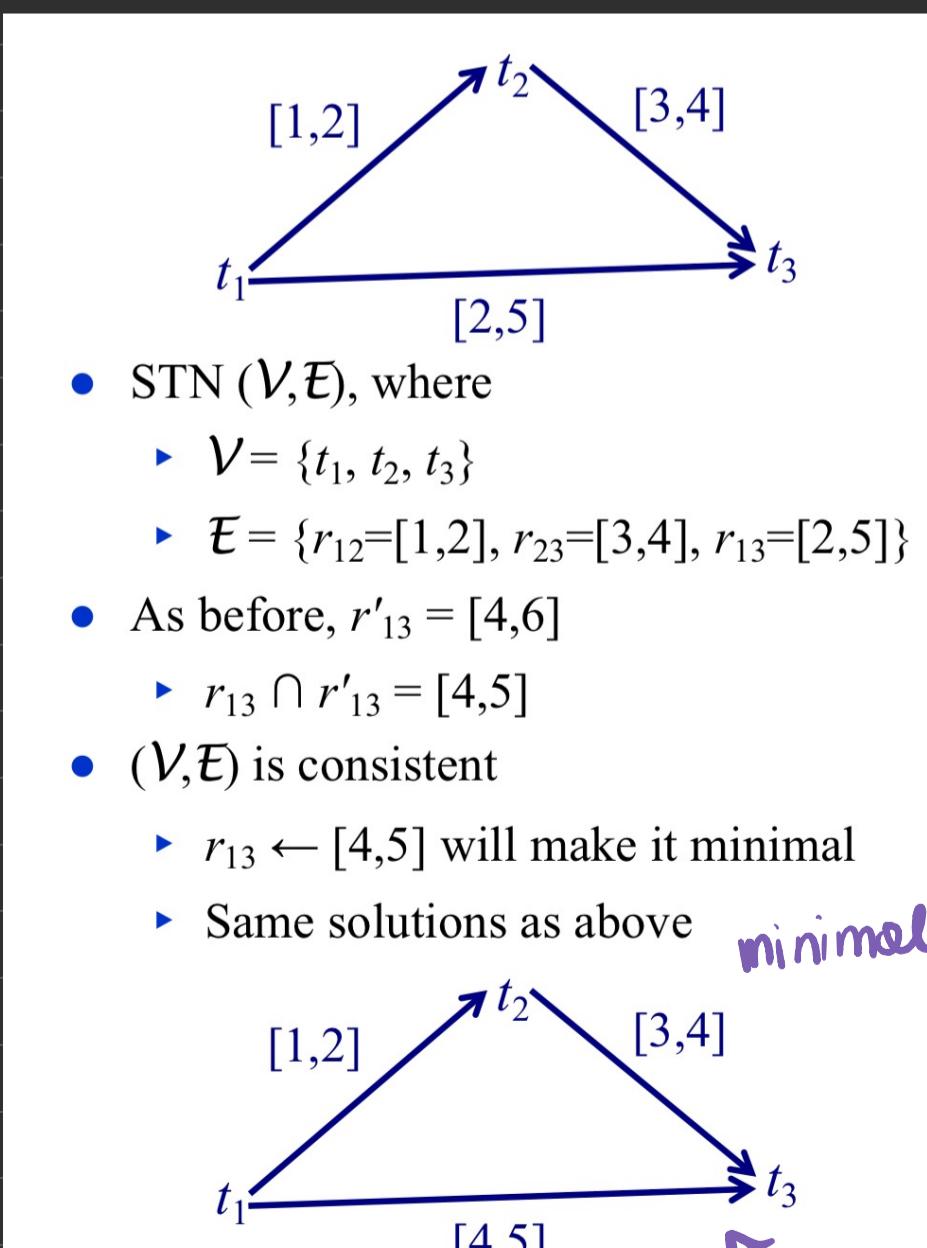
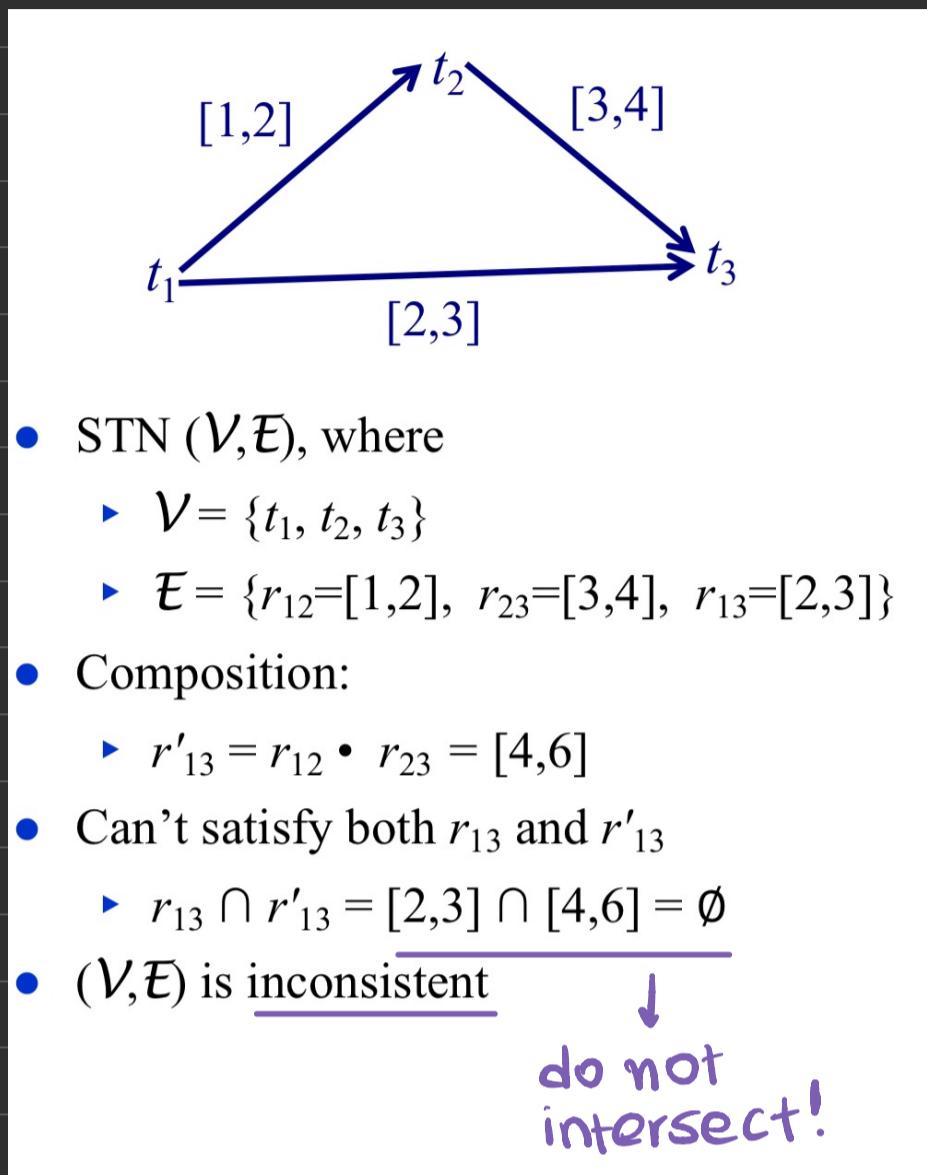


consistency checking

• intersection → finds common interval between 2 constraints

• composition → combines 2 constraints through a variable to form a direct constraint

examples:



$V = \{t_1, t_2, t_3\}$

$E = \{r_{12} = [1,2]$

$r_{23} = [3,4]$

$r_{13} = [2,5]\}$

composition :

$r'_{13} = [1+3, 2+4]$

$= [4, 6]$

$r'_{13} \cap r_{13} =$

$[4,6] \cap [2,5] =$

$[4,5] \neq \emptyset$

\downarrow
max min

$\Rightarrow (V, E)$ consistent

$r_{13} \leftarrow [4,5]$ will make it minimal

PATH CONSISTENCY ALGORITHM

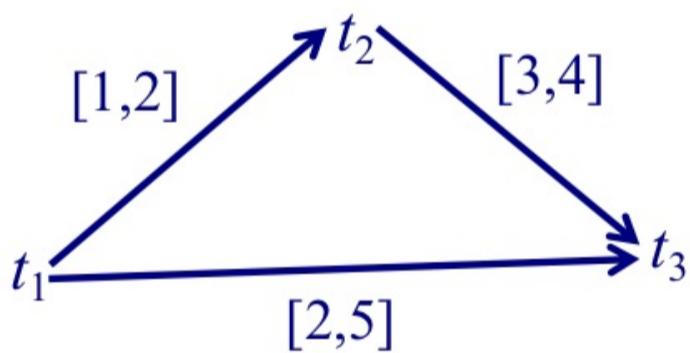
(PC)

- PC (*Path Consistency*) algorithm:

- ▶ For every $i < j$, iterates over r_{ij} once for each k
- ▶ Each time, tries to reduce interval, checks consistency
- ▶ i, j, k each go from 1 to n
=> time $O(n^3)$

$\mathcal{O}(n^3)$

Operations on STNs

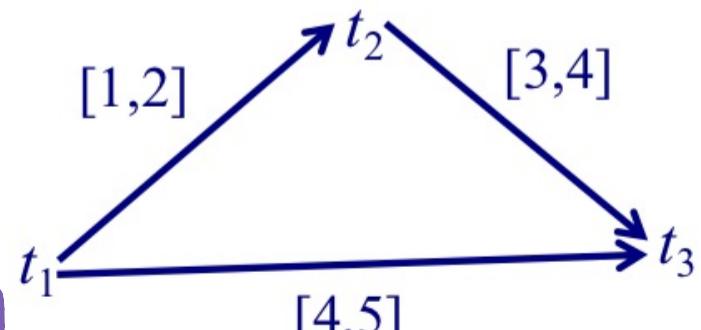


- For $n = 3$, PC tries these triples:

- ▶ $k=1, i=2, j=3$
- ▶ $k=2, i=1, j=3$
- ▶ $k=3, i=1, j=2$
- ▶ $k=1, i=2, j=3$
 - ▶ $r_{ik} = r_{21} = [-2, -1]$
 - ▶ $r_{kj} = r_{13} = [2, 5]$
 - ▶ $r_{ij} = r_{23} = [3, 4]$
 - ▶ $r_{ik} \cdot r_{kj} = [-2+2, -1+5] = [0, 4]$
 - ▶ $r_{23} = r_{ij} \leftarrow [\max(3, 0), \min(4, 4)] = [3, 4]$
 - ▶ No change $[0, 4] \cap [3, 4] = [3, 4]$

- $k=2, i=1, j=3$ (see previous slide)
 - ▶ PC reduces $r_{ij} = r_{13}$ to $[4, 5]$
- $k=3, i=1, j=2$
 - ▶ $r_{ik} = r_{13} = [4, 5]$
 - ▶ $r_{kj} = r_{32} = [-4, -3]$
 - ▶ $r_{ij} = r_{12} = [1, 2]$
 - ▶ $r_{ik} \cdot r_{kj} = [4-4, 5-3] = [0, 2]$
 - ▶ $r_{13} = r_{ij} \leftarrow [\max(1, -2), \min(2, 2)] = [1, 2]$ $[0, 2] \cap [1, 2] = [1, 2]$
 - ▶ No change $[1, 2]$

- Minimal network:



- The PC algorithm iteratively refines the constraints in the network by considering all possible paths.
- The process ensures that all constraints are consistent with each other, resulting in a minimal network where all intervals are as tight as possible without introducing inconsistency.
- This helps in verifying whether a given STN is consistent and if not, adjusting it to achieve consistency.